

Finding Heterogeneous-Unsafe Configuration Parameters in Cloud Systems

Sixiang Ma, Fang Zhou, Michael D. Bond, Yang Wang
The Ohio State University

EuroSys 2021



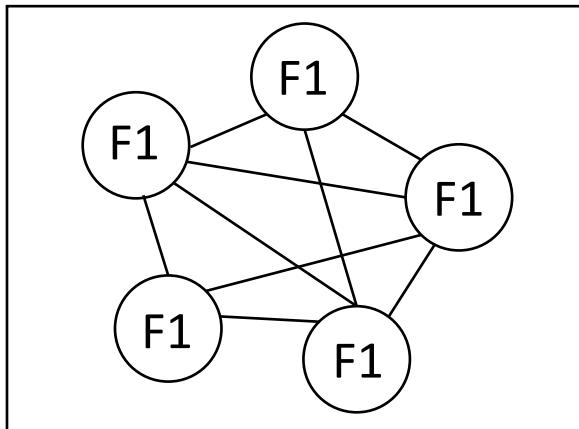
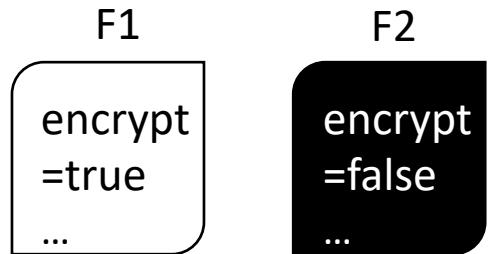
Heterogeneous Configurations Are Prevalent

- Heterogeneous hardware calls for heterogeneous configuration
- Online reconfiguration, e.g., reconfig command, rolling restart
 - Consequence: short window of heterogeneous configuration

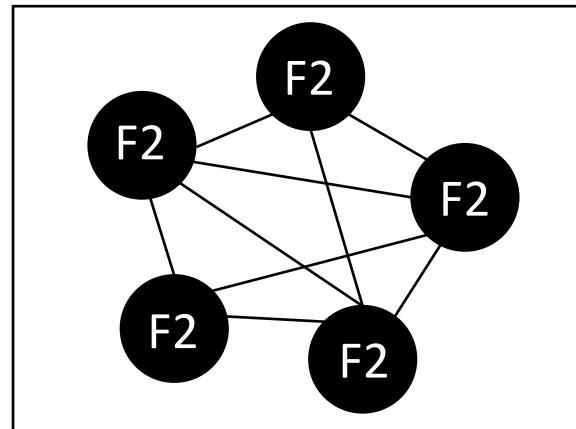


Heterogeneous Configuration Can Cause Error

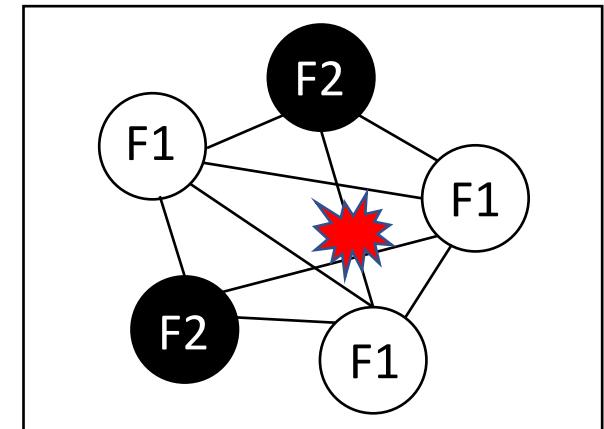
- Errors can happen even if each node has valid configuration locally.



HomoConf(F1) is valid



HomoConf(F2) is valid



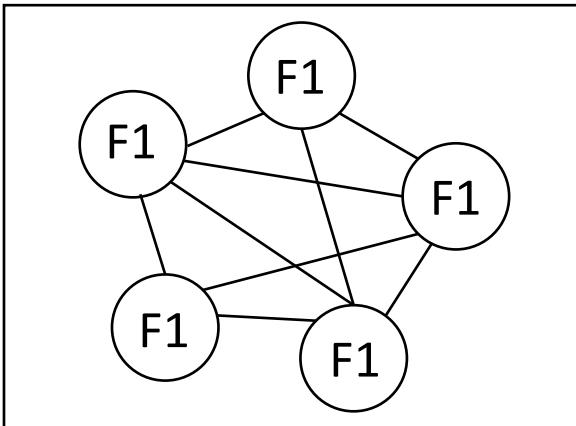
HeterConf(F1,F2) is invalid

Heterogeneous Configuration Can Cause Error

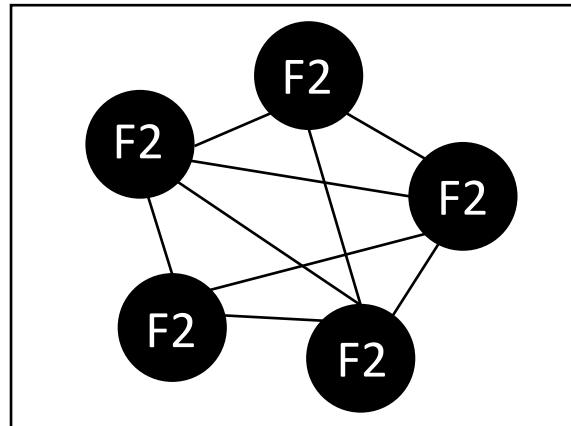
- Errors occur when F1 and F2 are heterogeneous.
- We call $\text{HeterConf}(F1, F2)$ ***Invalid Heterogenous Configuration***, if it causes errors but $\text{HomoConf}(F1)$ and $\text{HomoConf}(F2)$ do not.

F2

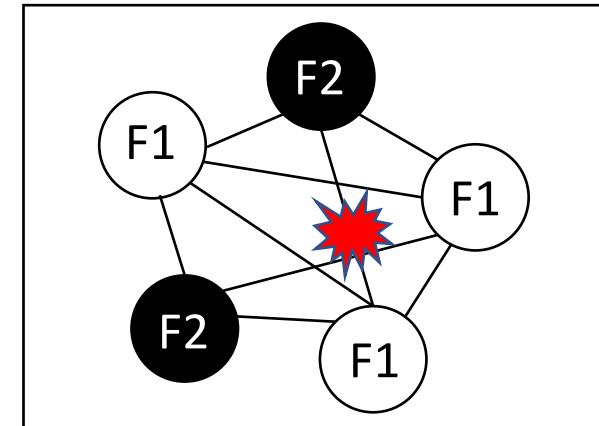
encrypt
=false
...



$\text{HomoConf}(F1)$ is valid



$\text{HomoConf}(F2)$ is valid



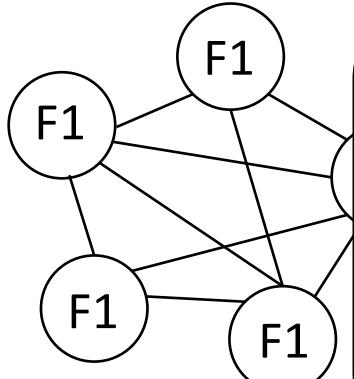
$\text{HeterConf}(F1, F2)$ is invalid

Heterogeneous Configuration Can Cause Error

- Errors valid We call HeterConf(F_1, F_2) ***Invalid Heterogenous Configuration***, if it causes errors but $\text{HomoConf}(F_1)$ and $\text{HomoConf}(F_2)$ do not.

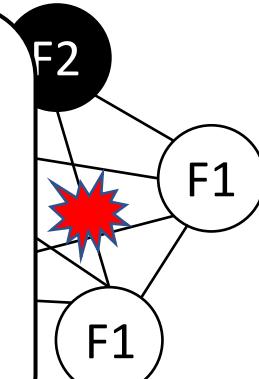
F_2

encrypt
=false
...



$\text{HomoConf}(F_1)$ is valid

We call the corresponding parameter
Heterogenous-Unsafe Configuration Parameter.



$\text{HomoConf}(F_2)$ is valid

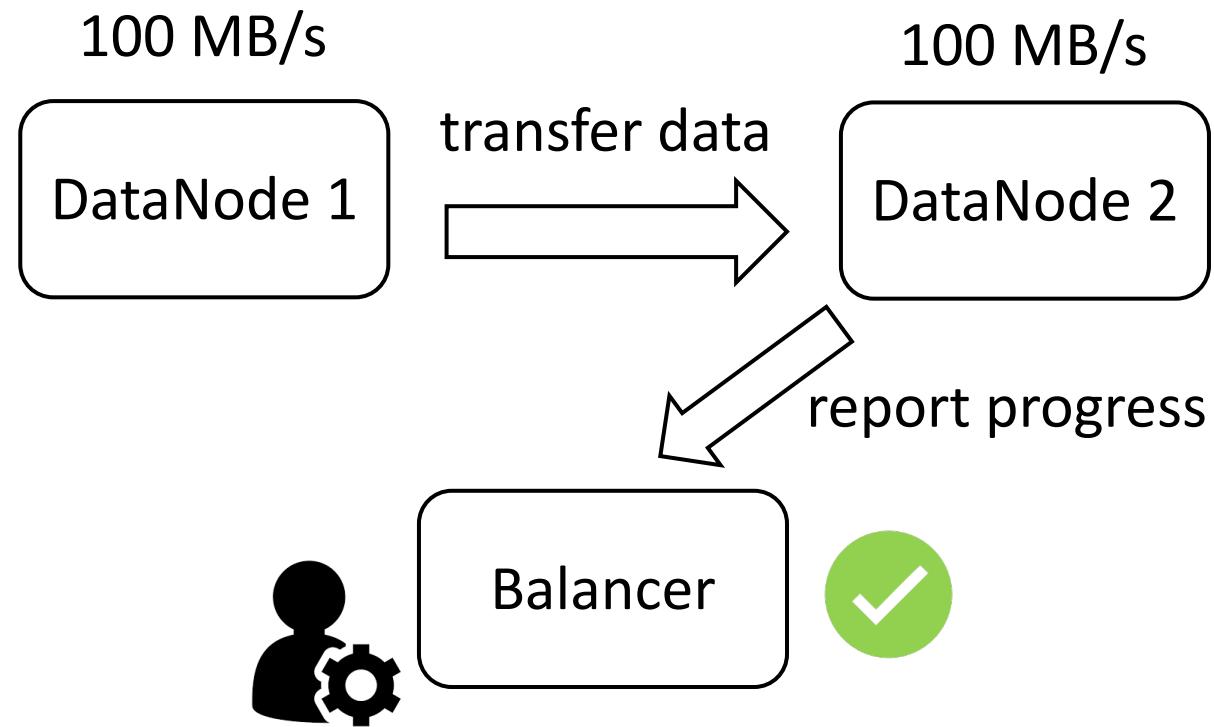
$\text{HeterConf}(F_1, F_2)$ is invalid

Example: *dfs.datanode.balance.bandwidthPerSec*

- Specify the maximum amount of bandwidth that a HDFS DataNode can use for balancing purpose.

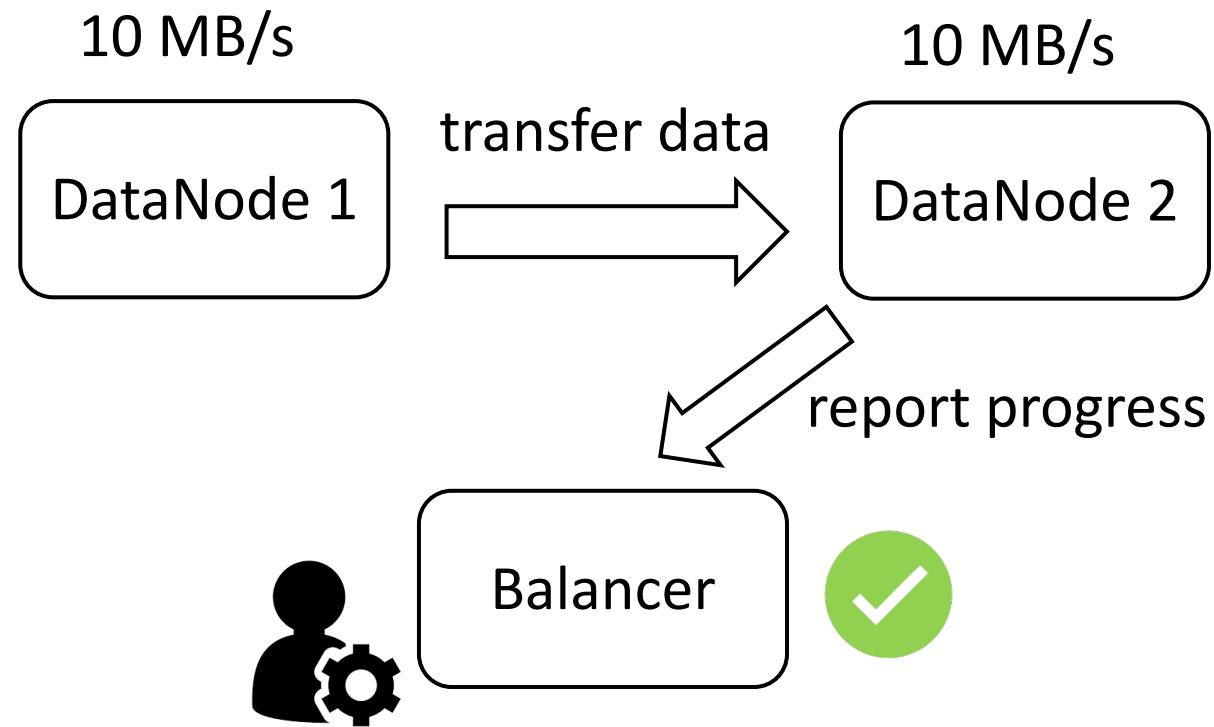
Example: *dfs.datanode.balance.bandwidthPerSec*

- Specify the maximum amount of bandwidth that a HDFS DataNode can use for balancing purpose.



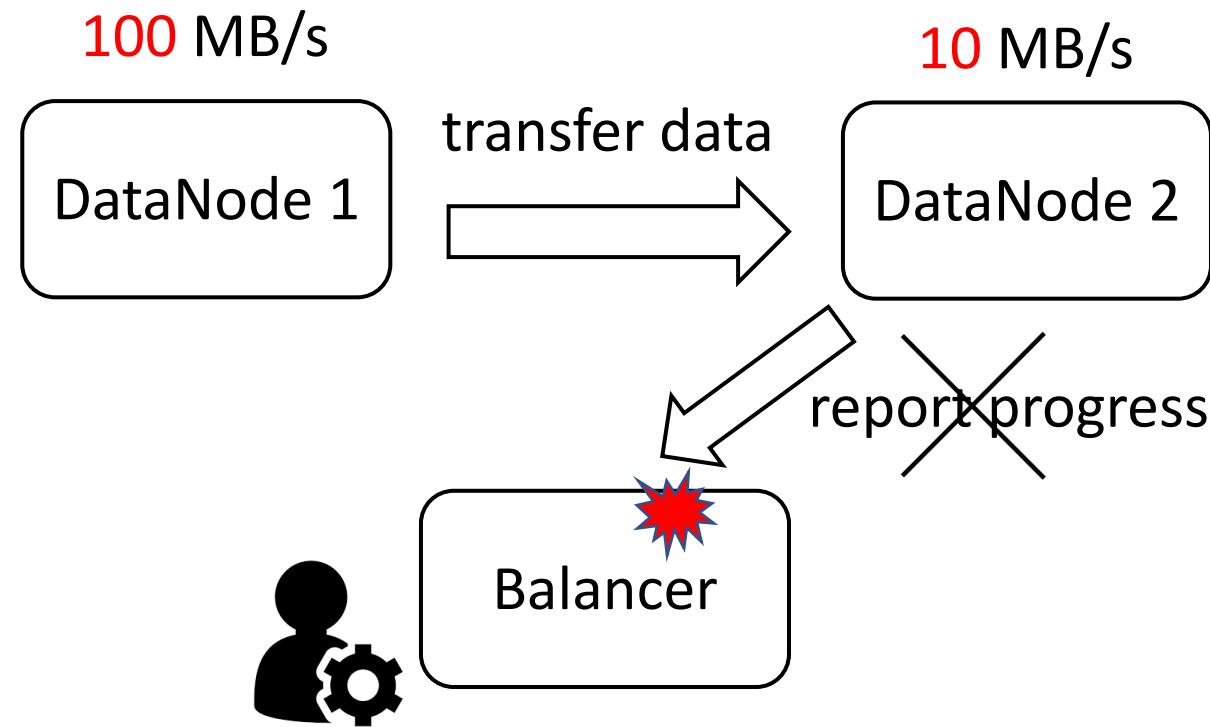
Example: *dfs.datanode.balance.bandwidthPerSec*

- Specify the maximum amount of bandwidth that a HDFS DataNode can use for balancing purpose.



Example: *dfs.datanode.balance.bandwidthPerSec*

- Specify the maximum amount of bandwidth that a HDFS DataNode can use for balancing purpose.



Related Work

This type of errors is different from the problem of erroneous configuration values [EnCore-ASPLOS'14, ConfValley-EuroSys'15, PCheck-OSDI'16, PracExtractor-ATC'20]

- Parameter values are valid.
- Errors happen when nodes communicate.

Overview

- Our goal: find heterogeneous-unsafe configuration parameters in cloud systems.
- ZebraConf: a testing framework that reuse existing unit tests
- It finds 41 true problems in HDFS, YARN, MR, HBase, Flink.



ZebraConf Uses Classic Software Testing Approach

- Challenge: some parameters may only take effect under specific workloads.
- Observation: mature cloud systems usually have rich unit tests.
 - High code coverage [Kairux-SOSP'19]
 - E.g., 90.1% statement coverage in HDFS
 - Many unit tests are using configuration
 - 3,628 unit tests in HDFS use config, covering 96.2% parameters

ZebraConf Uses Classic Software Testing Approach

- Challenge: some parameters may only take effect under specific workloads.

Reuse Existing Unit Tests for Our Purpose

- E.g., 90.1% statement coverage in HDFS
- Many unit tests are using configuration
 - 3,628 unit tests in HDFS use config, covering 96.2% parameters

ZebraConf: Major Challenges

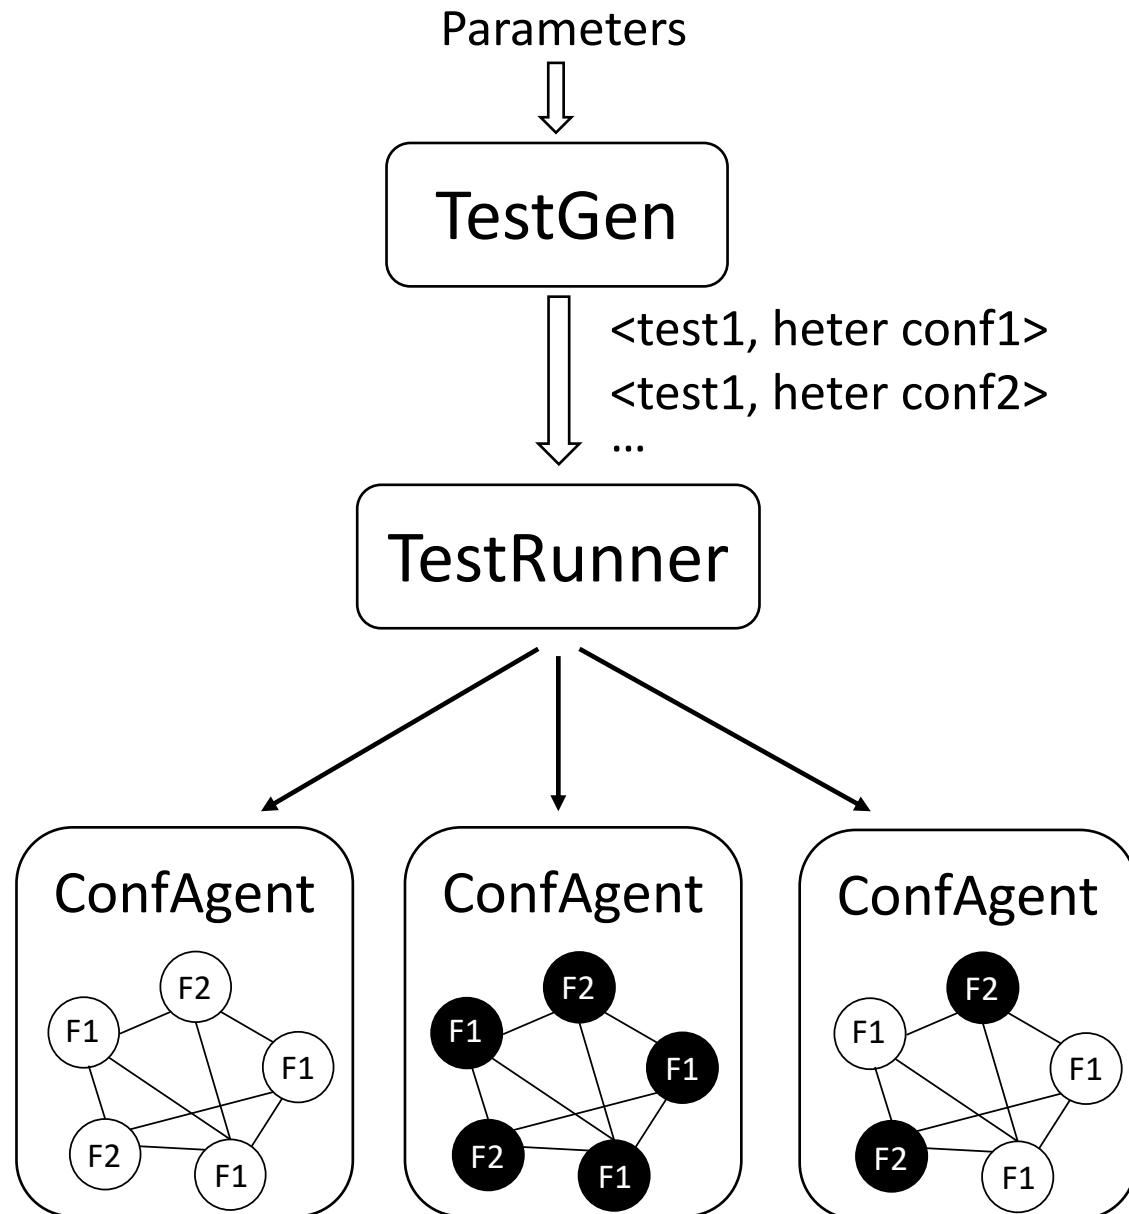
ZebraConf: Major Challenges

- C1: How to reduce testing time?
 - Apps can have 1000s of tests, 100s-1000s of parameters.
 - A test runs for several seconds to several minutes.

ZebraConf: Major Challenges

- C1: How to reduce testing time?
 - Apps can have 1000s of tests, 100s-1000s of parameters.
 - A test runs for several seconds to several minutes.
- C2: How to assign heterogeneous configuration in unit tests?
 - We can specify the config when starting a node as process.
 - E.g., `hadoop-daemon.sh --config [CONFIG_PATH] start`
 - However, this approach doesn't work in unit tests.

ZebraConf Overview

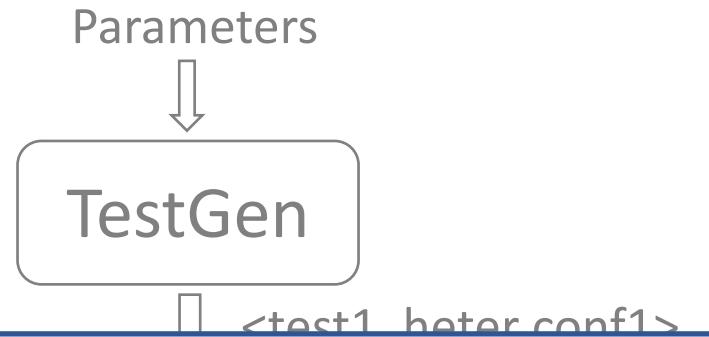


TestGen: generate test instances

TestRunner: conduct a test

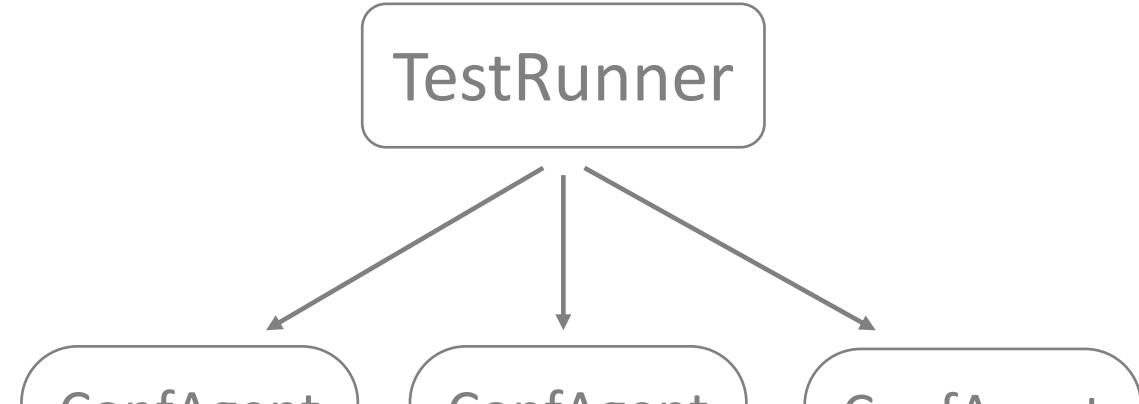
ConfAgent: assign configs to nodes

ZebraConf Overview



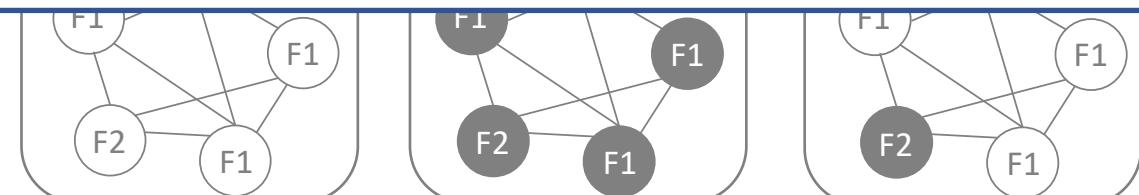
TestGen: generate test instances

C1: How to reduce testing time



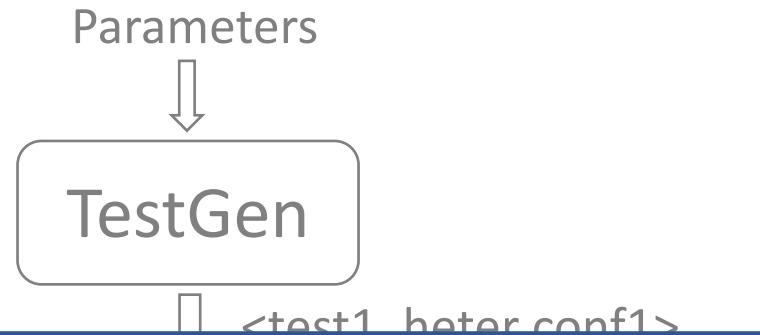
TestRunner: conduct a test

C2: How to assign heter config in unit tests

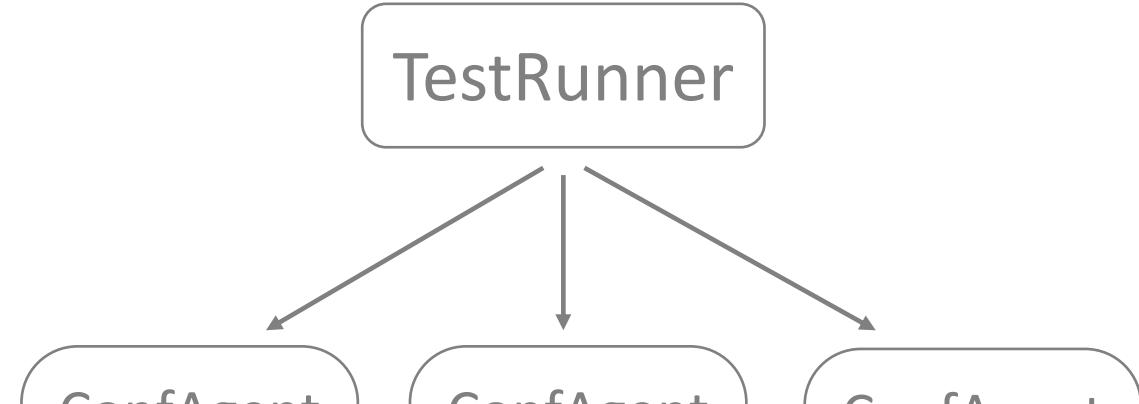


ConfAgent: assign configs to nodes

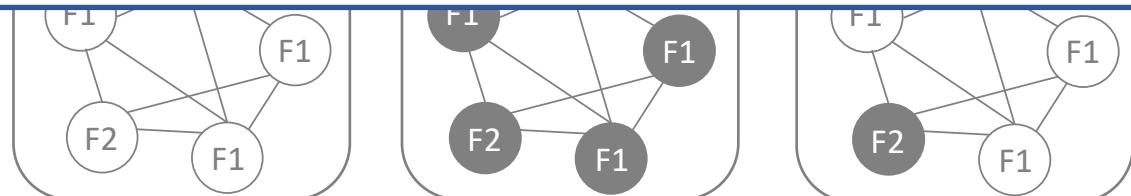
ZebraConf Overview



C1: How to reduce testing time



C2: How to assign heter config in unit tests



TestGen: generate test instances

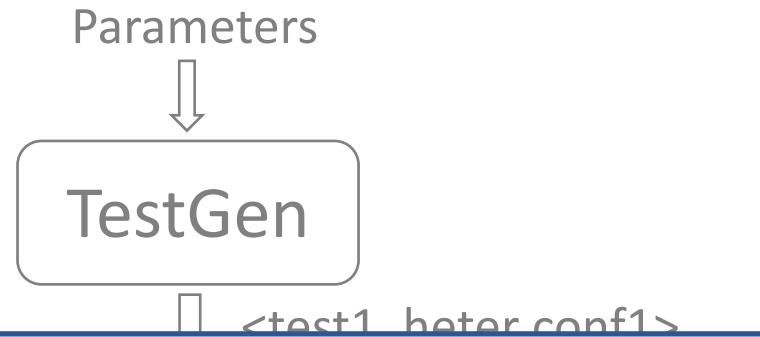
- Selective value assignment
- Pre-run profiling
- Pooled tests

TestRunner: conduct a test

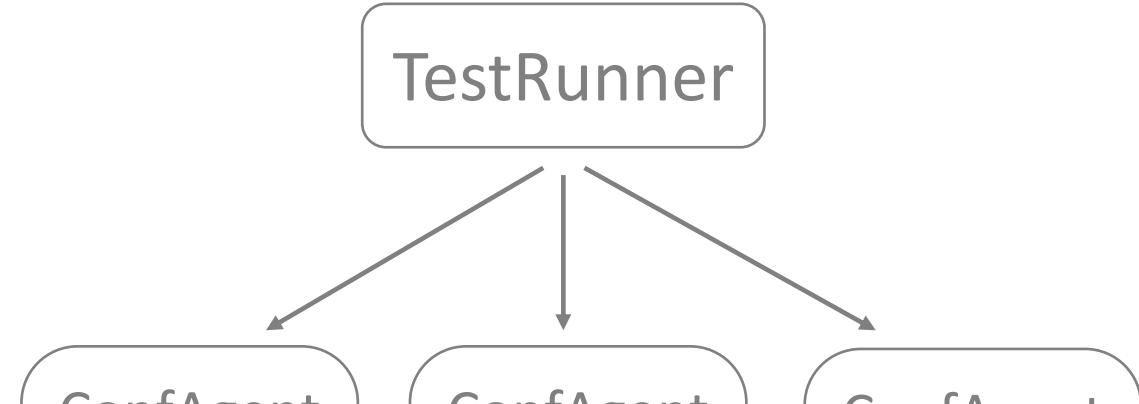
- Supporting pooled testing
- Concurrent testing
- Hypothesis testing

ConfAgent: assign configs to nodes

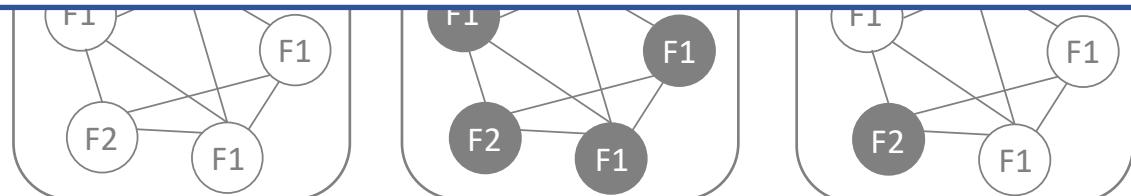
ZebraConf Overview



C1: How to reduce testing time



C2: How to assign heter config in unit tests



TestGen: generate test instances

- Selective value assignment
- Pre-run profiling
- Pooled tests

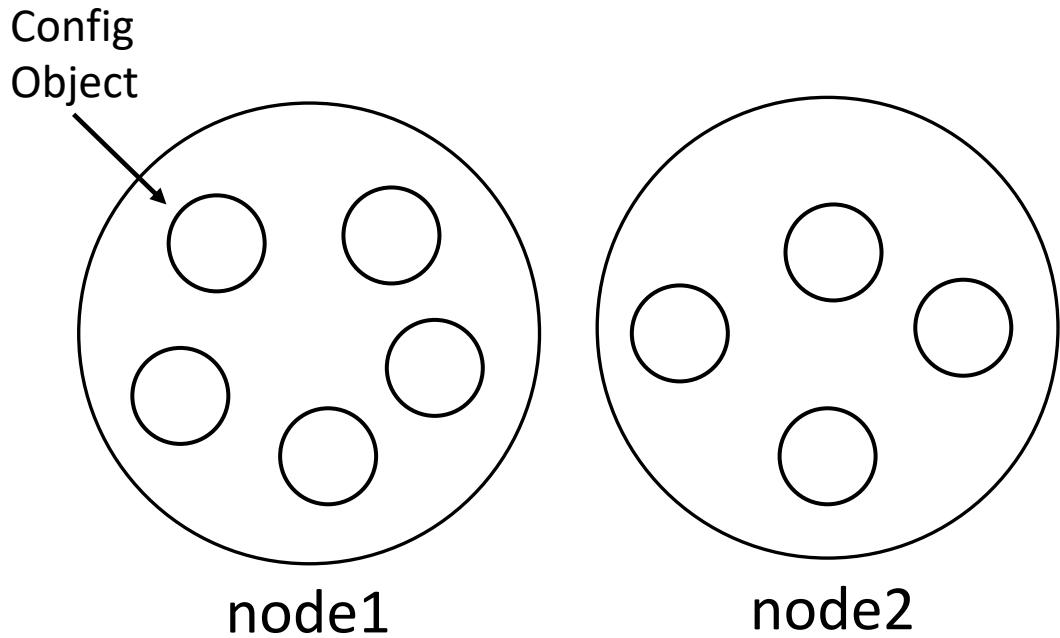
TestRunner: conduct a test

- Supporting pooled testing
- Concurrent testing
- Hypothesis testing

ConfAgent: assign configs to nodes

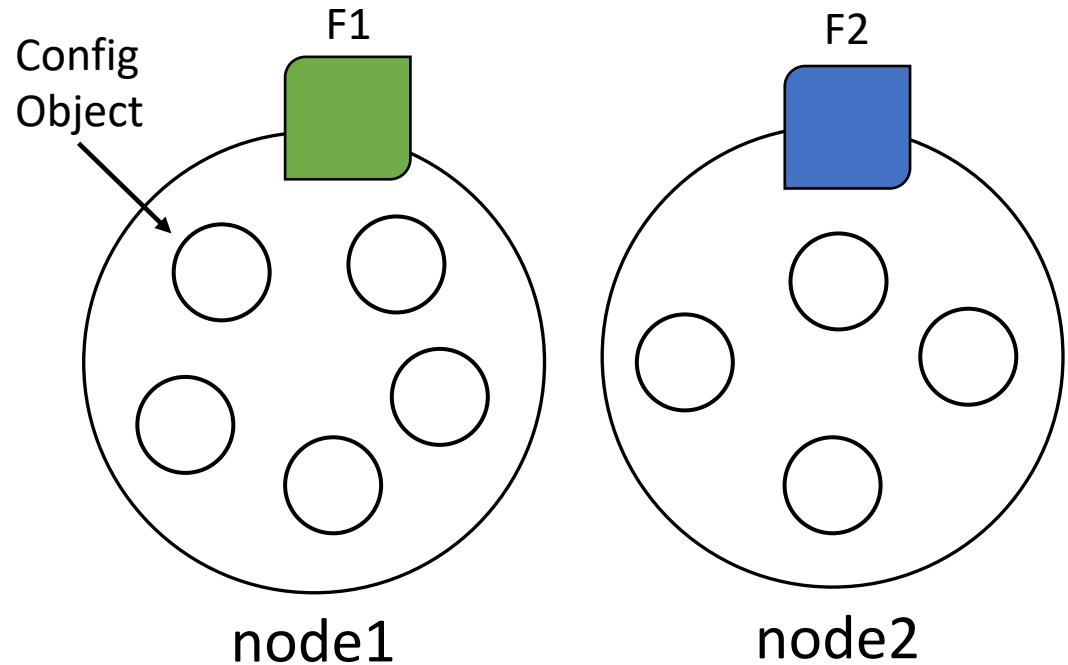
ConfAgent: Challenges

In distributed setting, we can specify the config file when starting a node as process.



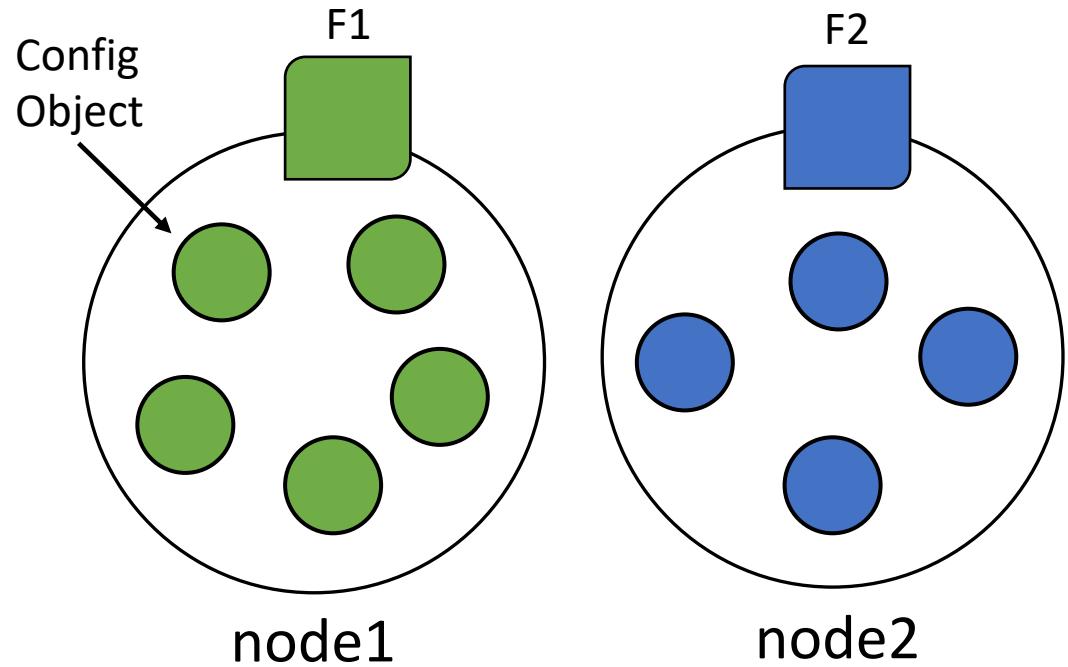
ConfAgent: Challenges

In distributed setting, we can specify the config file when starting a node as process.



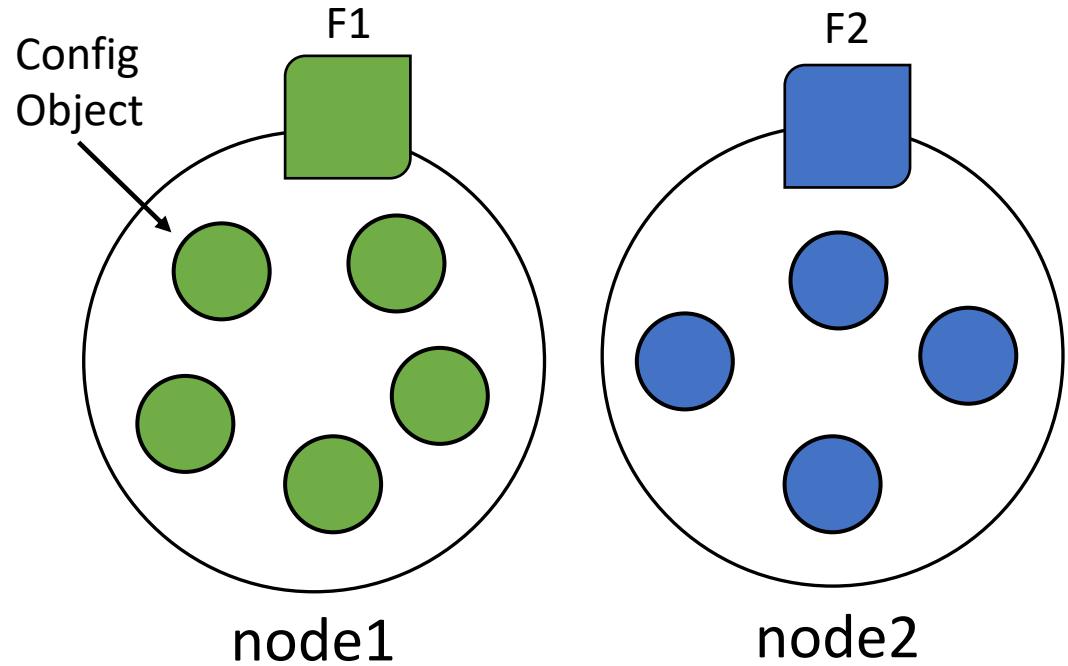
ConfAgent: Challenges

In distributed setting, we can specify the config file when starting a node as process.

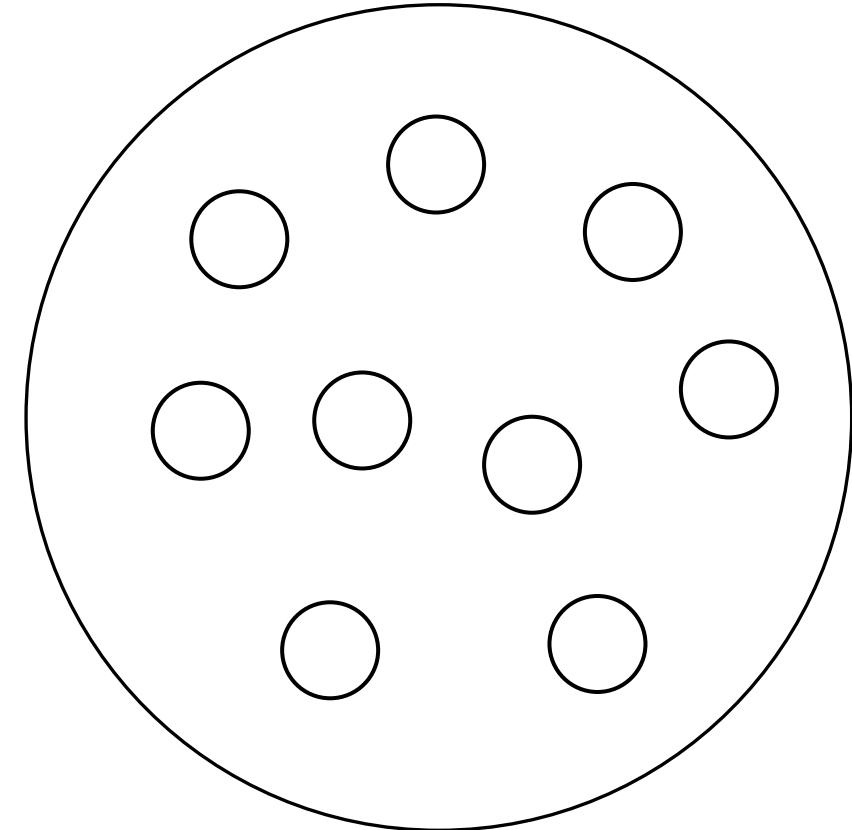


ConfAgent: Challenges

In distributed setting, we can specify the config file when starting a node as process.



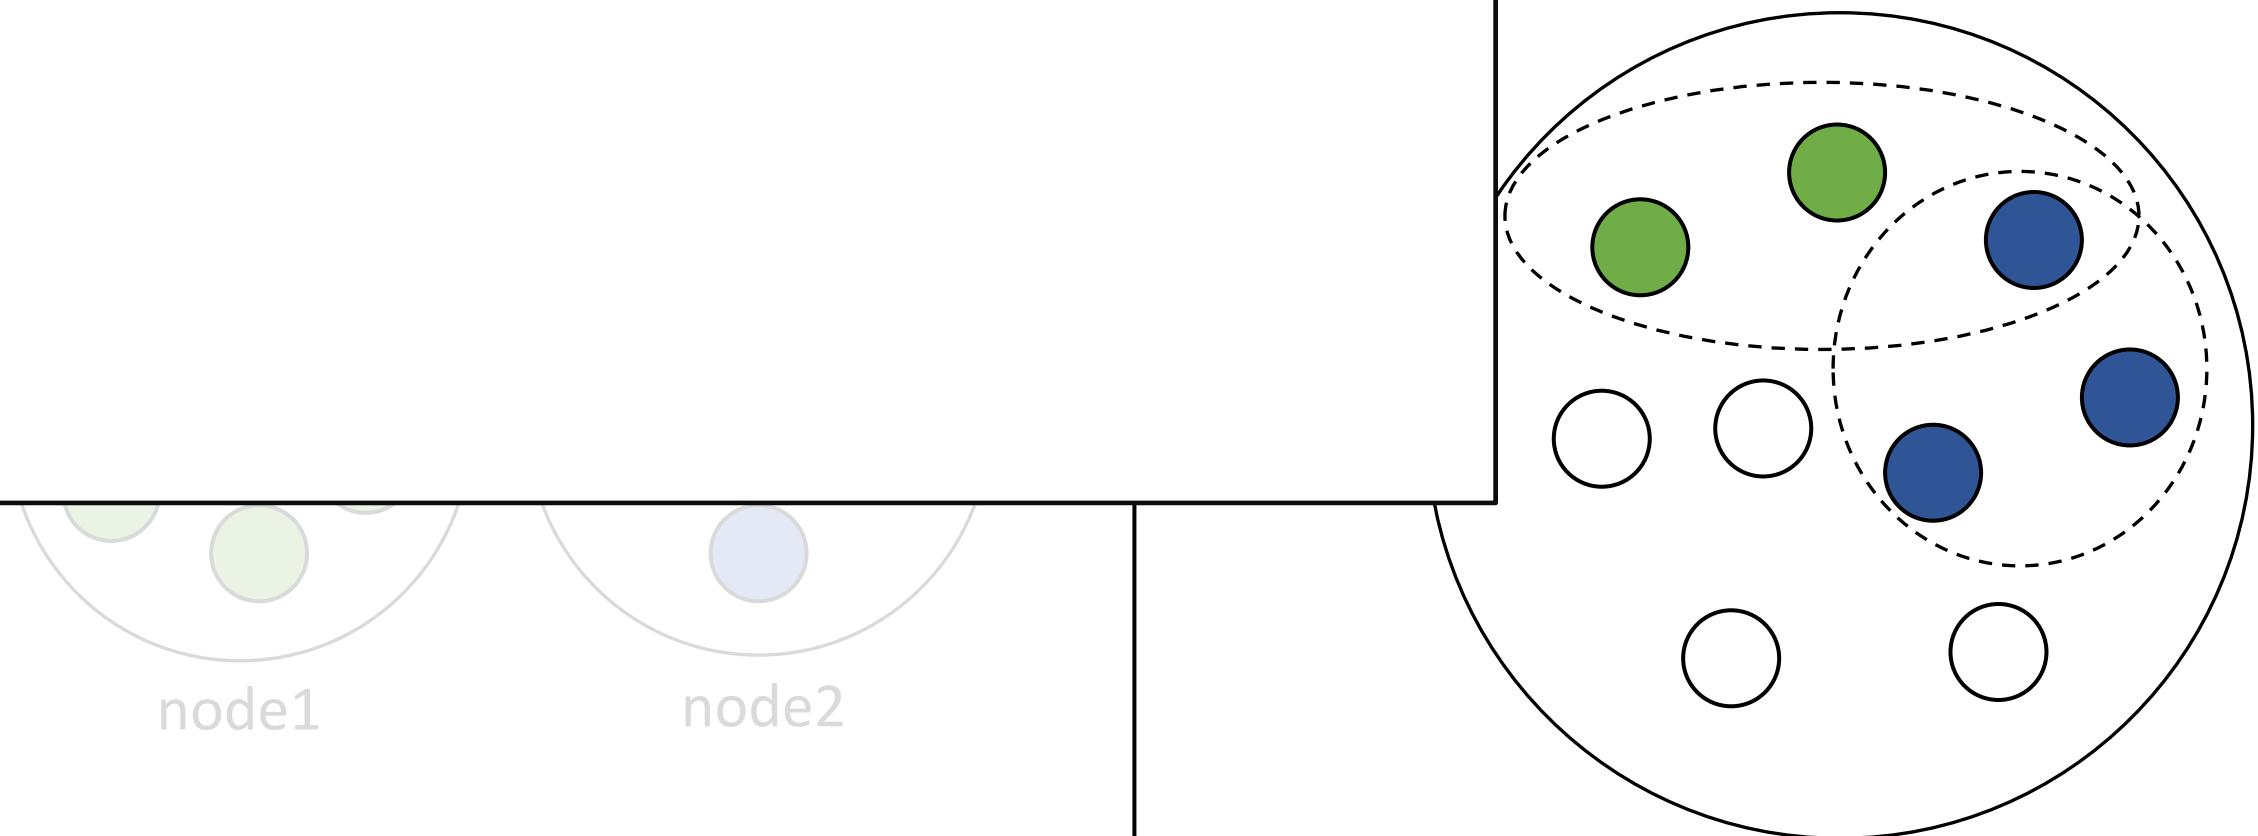
This doesn't work in unit tests, as nodes are often created as threads in a single process (i.e., minicluster testing)



ConfAgent: Challenges

Key idea: attribute config objects to nodes & hack return values

Can't work in unit tests, as nodes are created as threads in a single process (cluster testing)



ConfAgent: Challenges

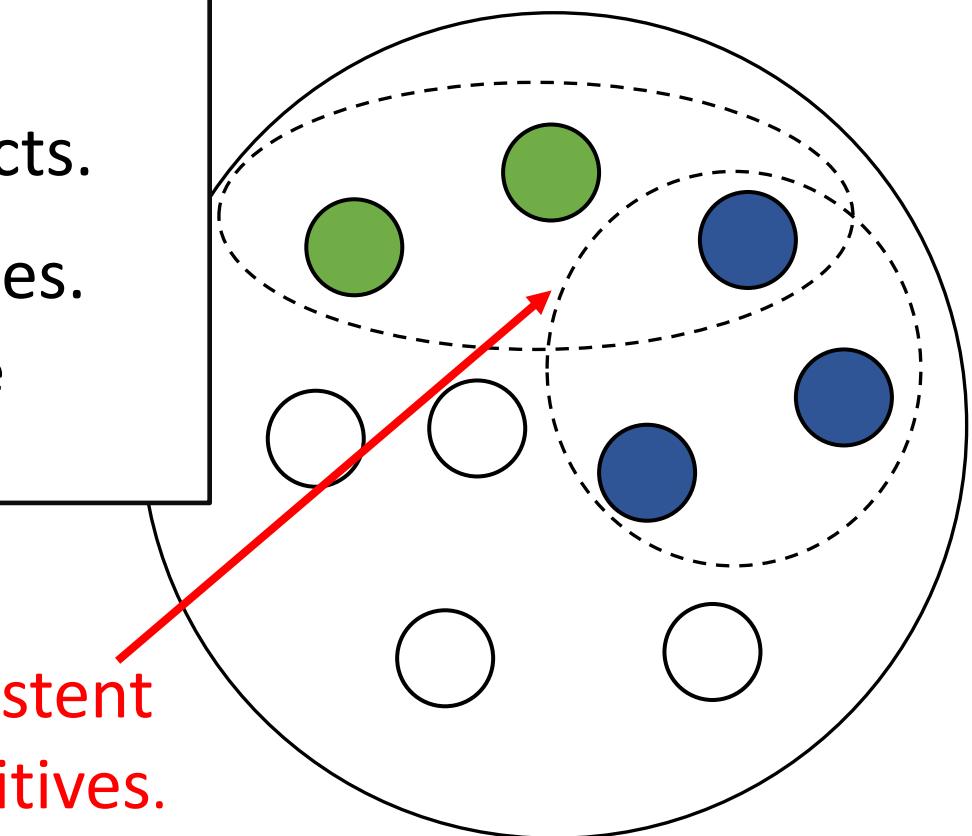
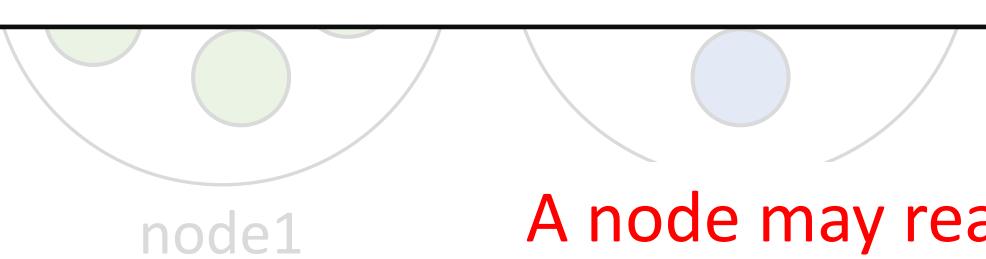
Key idea: attribute config objects to nodes & hack return values

Challenges:

- Each node can have multiple config objects.
- Config objects can be shared among nodes.
 - Values in one config object seen by multiple nodes.

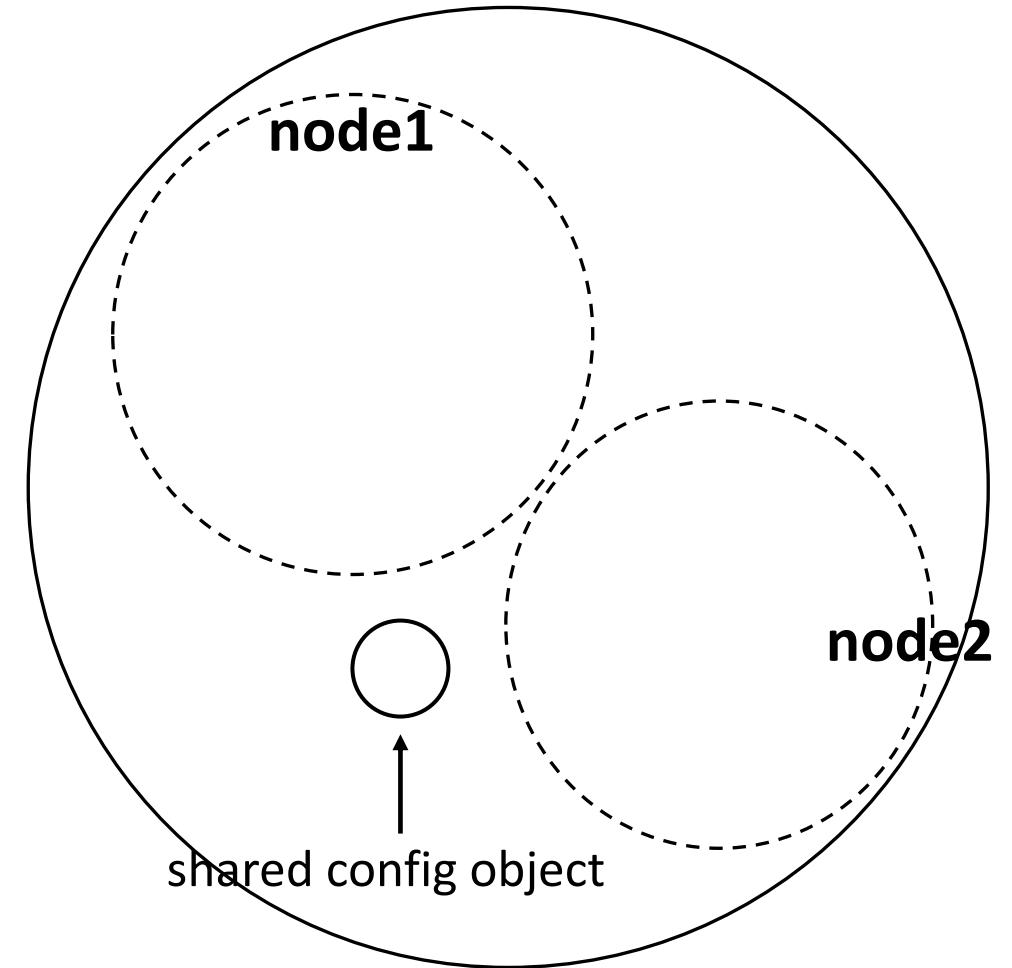
It work in unit tests, as nodes are used as threads in a single process (cluster testing)

A node may read inconsistent values, causing false positives.



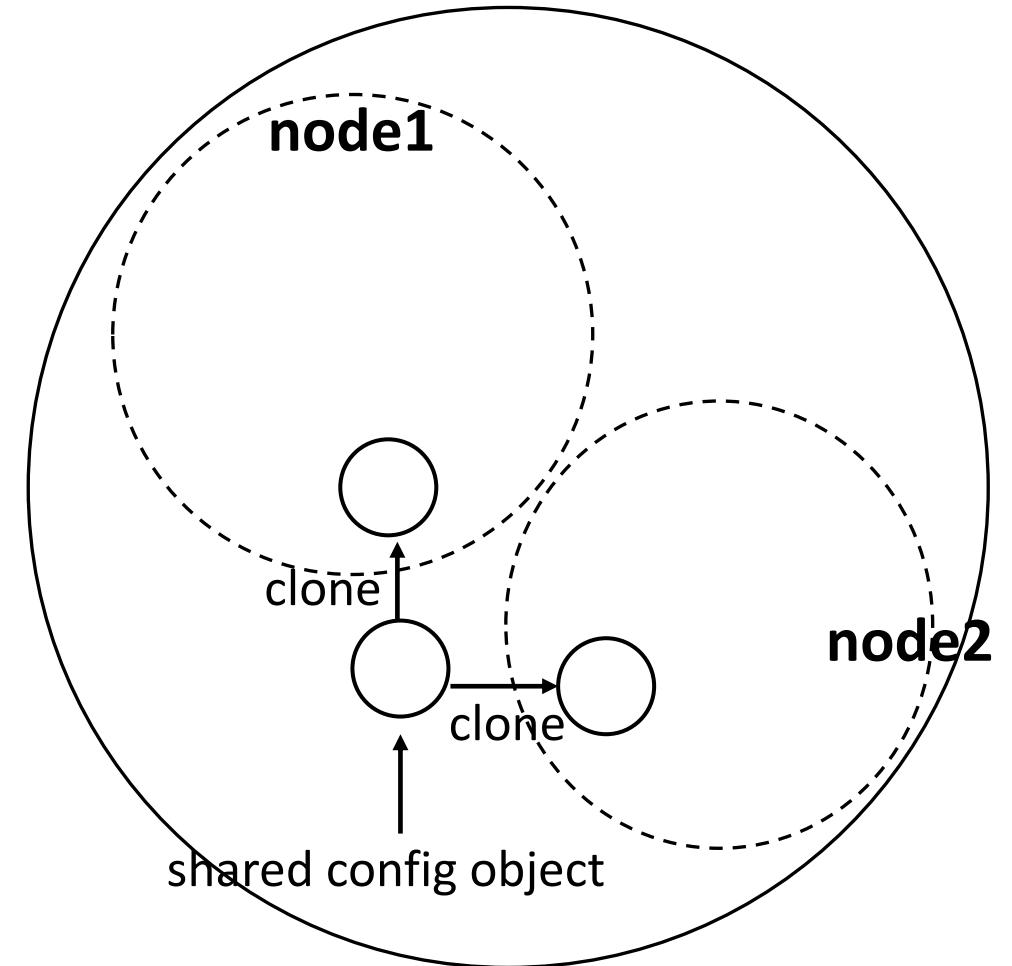
ConfAgent's Solution

- Clone config objects that can be shared.



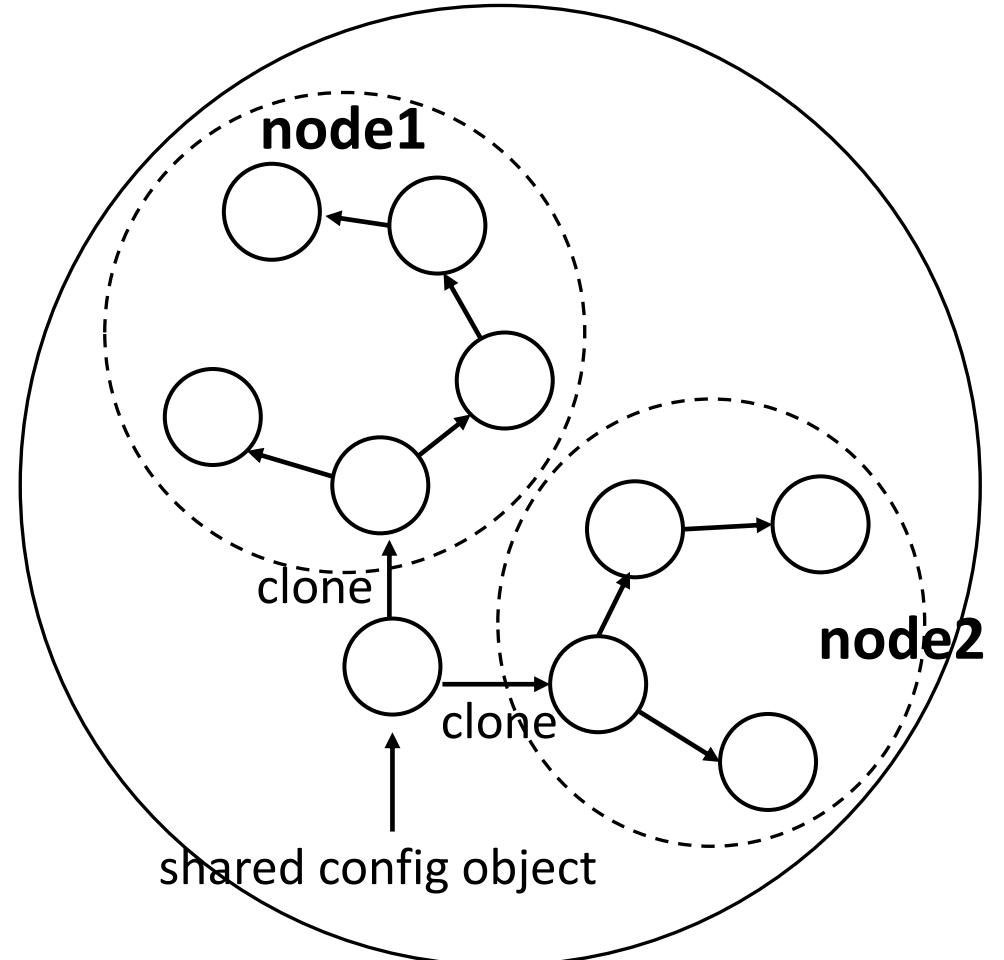
ConfAgent's Solution

- Clone config objects that can be shared.



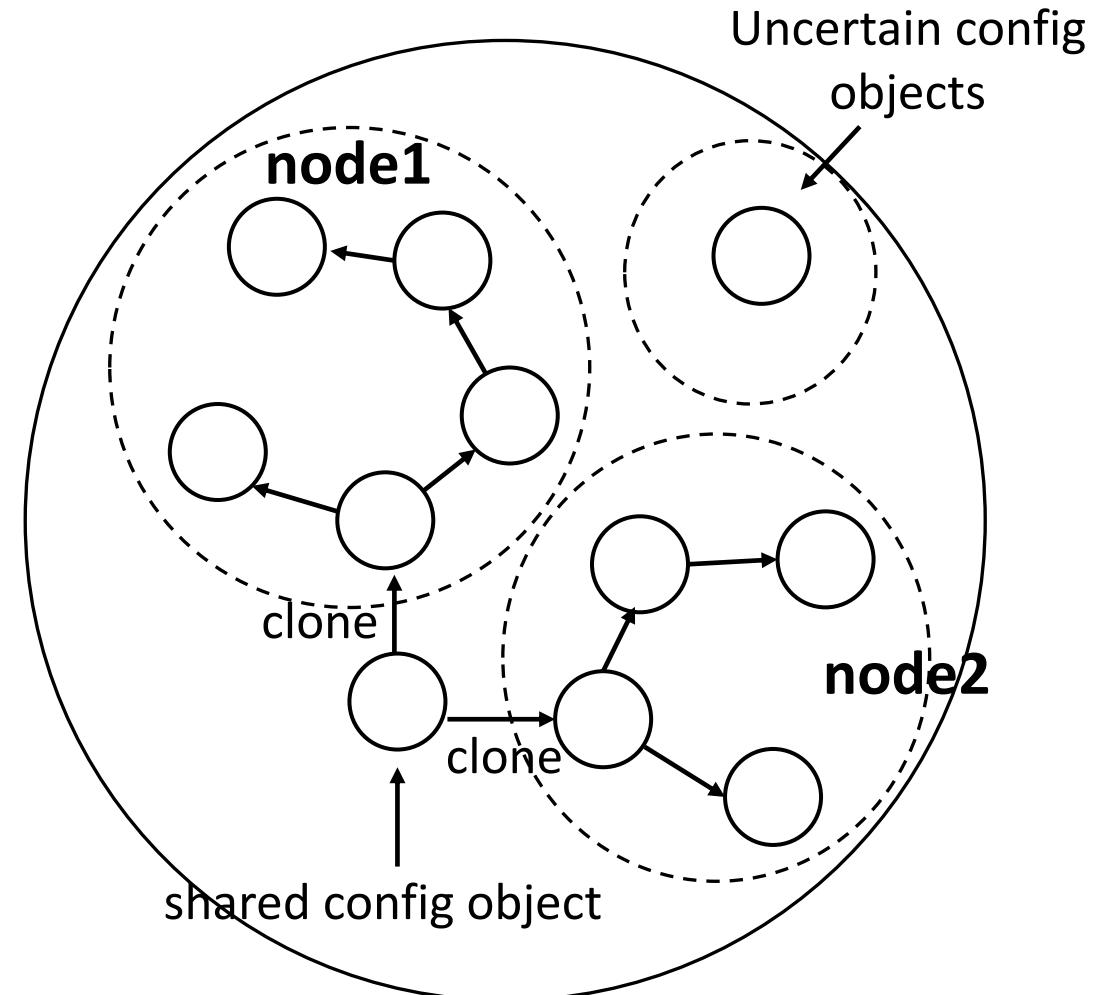
ConfAgent's Solution

- Clone config objects that can be shared.
- Track config creation flow and attribute config objects to nodes by rules.
 - E.g., Conf b = new Conf(a) & a belongs to node1
→ b belongs to node1



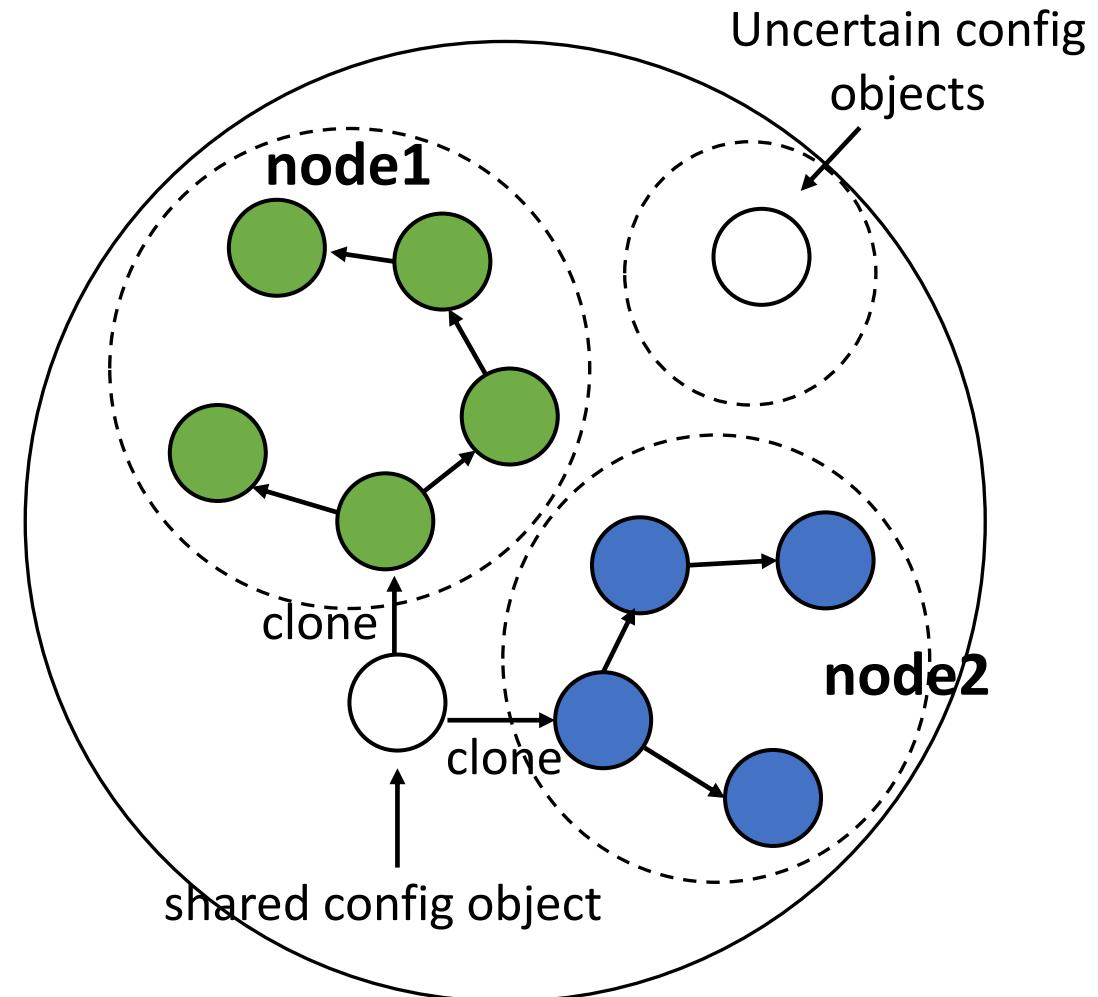
ConfAgent's Solution

- Clone config objects that can be shared.
- Track config creation flow and attribute config objects to nodes by rules.
 - E.g., Conf b = new Conf(a) & a belongs to node1
→ b belongs to node1
- Track uncertain config objects.
 - Avoid testing parameters read from them.



ConfAgent's Solution

- Clone config objects that can be shared.
- Track config creation flow and attribute config objects to nodes by rules.
 - E.g., Conf b = new Conf(a) & a belongs to node1
→ b belongs to node1
- Track uncertain config objects.
 - Avoid testing parameters read from them.
- Manipulate config parameter values.



Evaluation

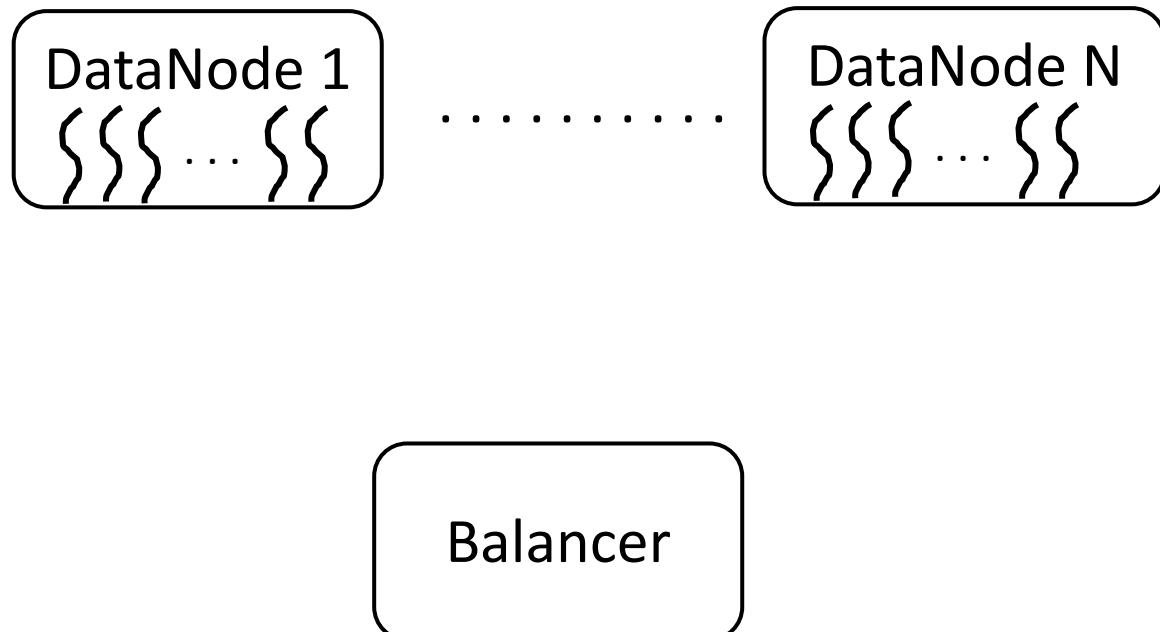
- Hardware setting
 - We run all the experiments on CloudLab
 - Intel Xeon 10-core CPUs, 192 GB DRAM, 480 GB SATA SSD
- Applications
 - Five app: HDFS, YARN, MR, HBase, Flink
 - Modification overhead: 18 to 38 LOC
 - Totally 4,652 machine hours with up to 100 physical machines, each running 20 Docker containers

Evaluation

- ZebraConf reports 57 heterogeneous-unsafe parameters.
- Our manual analysis finds 41 are true problems.
- Categories of these parameters:
 - Data transfer format related
 - Max limit related
 - Timeout related
 - Task numbers related
 - ...
 - Unexpected ones

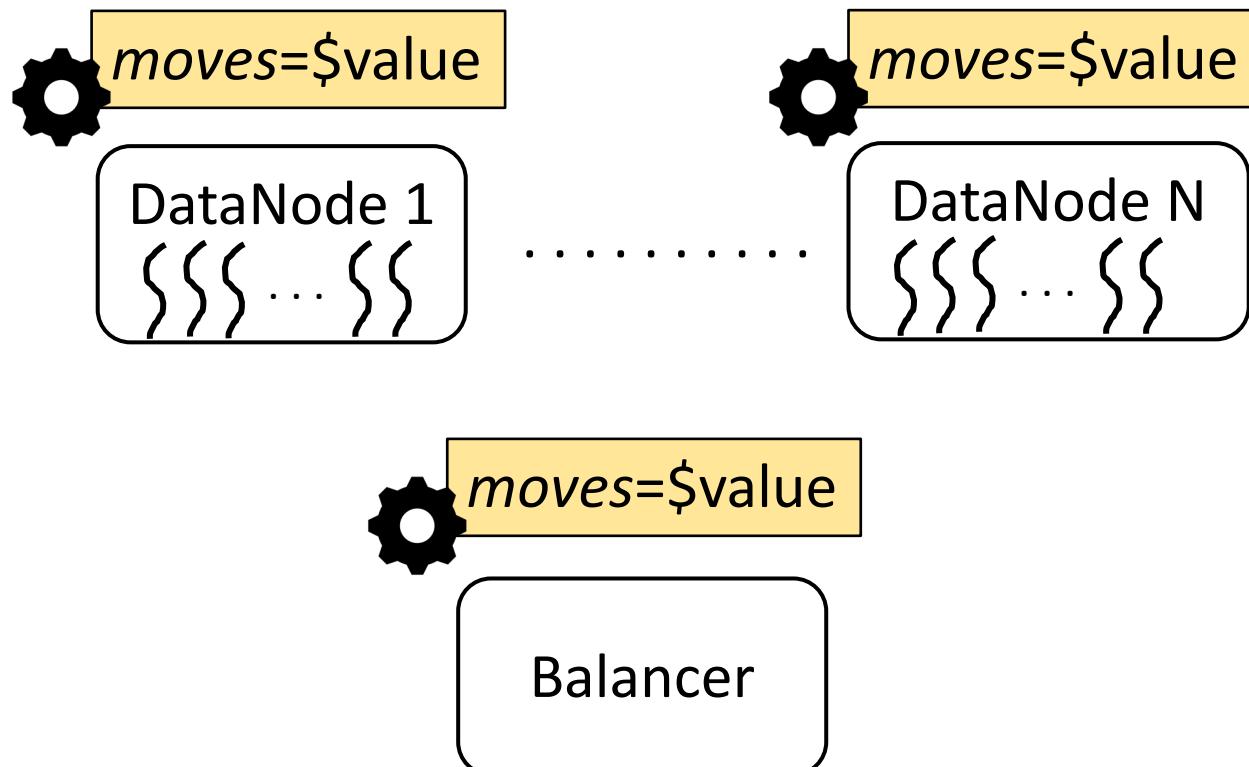
dfs.datanode.balance.max.concurrent.moves

- Limits the max number of threads that a DataNode can use for balancing.



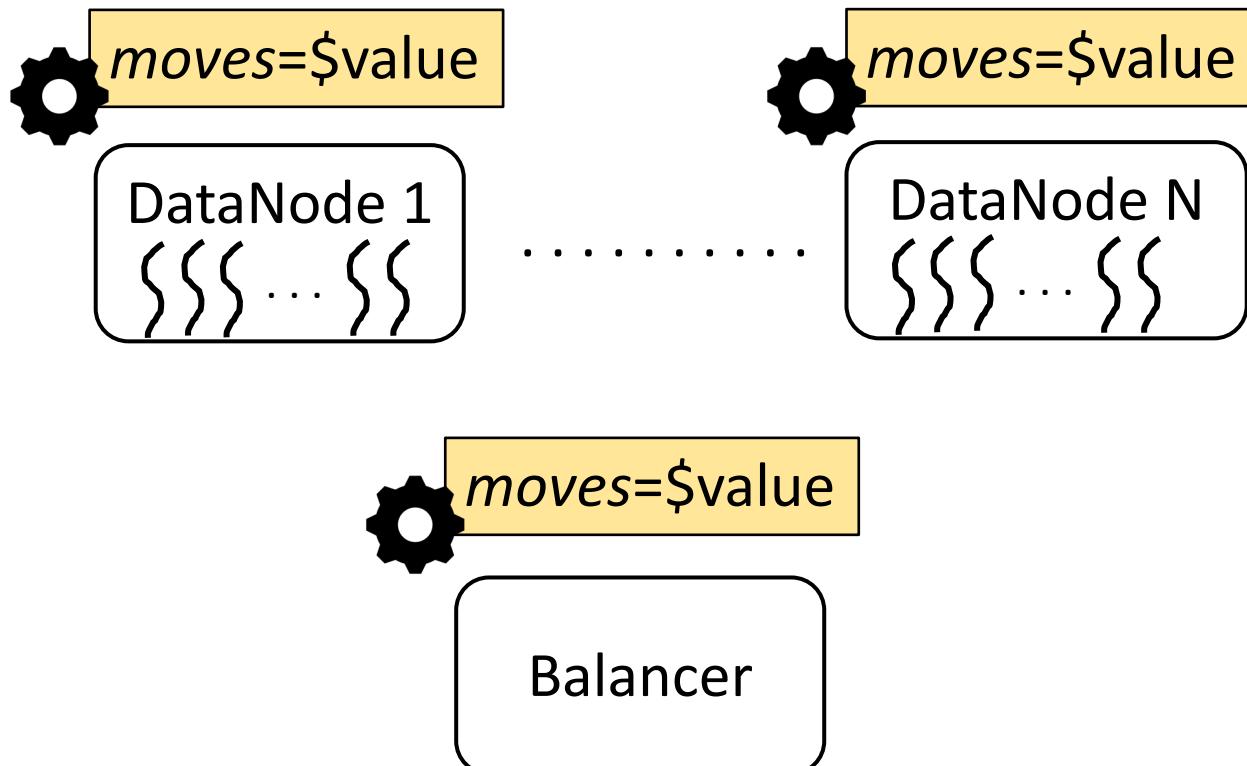
dfs.datanode.balance.max.concurrent.moves

- Limits the max number of threads that a DataNode can use for balancing.



dfs.datanode.balance.max.concurrent.moves

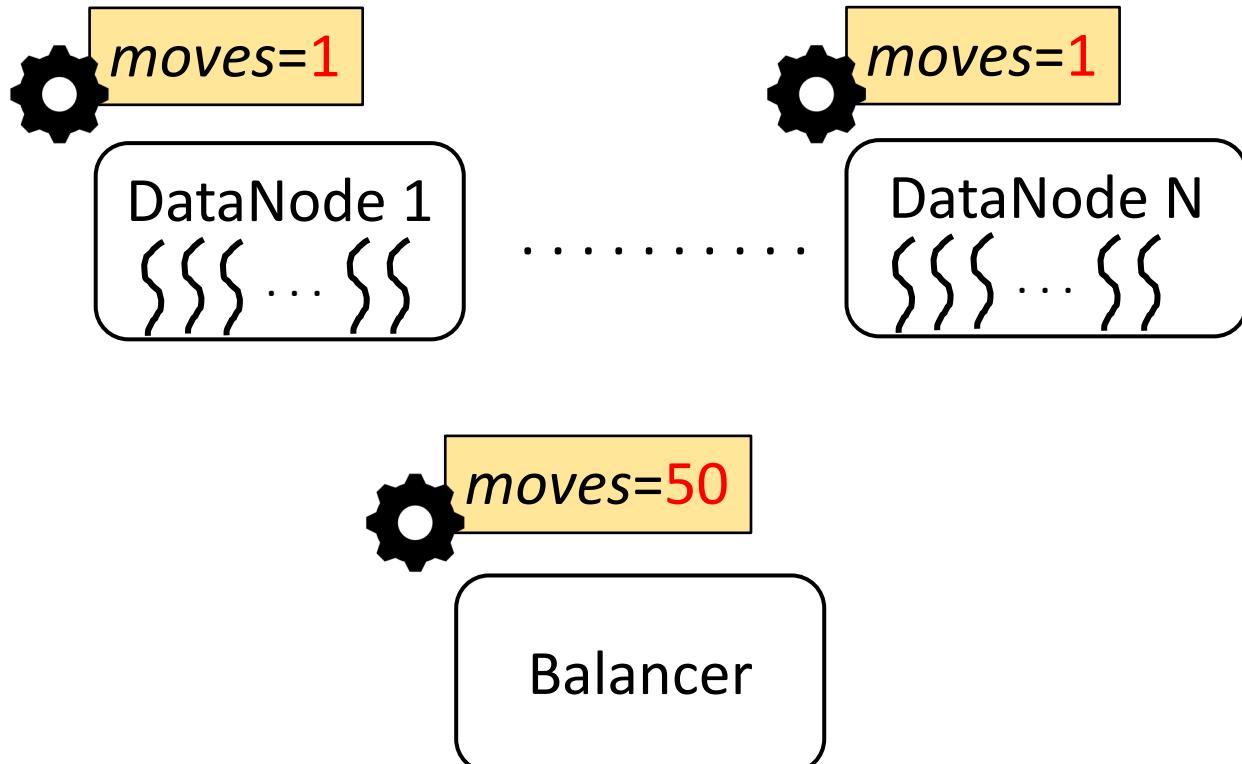
- Limits the max number of threads that a DataNode can use for balancing.



Configs	Balancing Time
Balancer: 50, DN: 50	14s
Balancer: 1, DN: 1	16.7s

dfs.datanode.balance.max.concurrent.moves

- Limits the max number of threads that a DataNode can use for balancing.

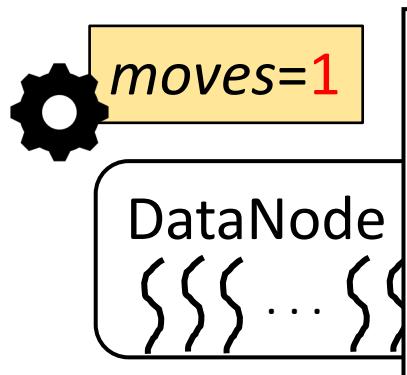


Configs	Balancing Time
Balancer: 50, DN: 50	14s
Balancer: 1, DN: 1	16.7s
Balancer: 50, DN: 1	154s

10x slower than just
using 1 thread

dfs.datanode.balance.max.concurrent.moves

- Limits the max number of threads that a DataNode can use for balancing.



HDFS-7466: Allow different values for *moves* per DataNode

“The correct approach will be to obtain the value from the DataNode itself.”

Balancer

Balancing Time	
0	14s
	16.7s
	154s

10x slower than just using 1 thread

Conclusion

- ZebraConf reuses existing unit tests to find unsafe parameters.
- We find 41 heterogeneous-unsafe parameters with ZebraConf.
- Need better support for heterogeneous configurations.
- We made ZebraConf publicly available:
<https://github.com/StarThinking/ZebraConf/>