

# PaSh: Light-Touch Data-Parallel Shell Processing

Nikos Vasilakis\*  
MIT

Konstantinos Kallas\*  
University of Pennsylvania

Konstantinos Mamouras  
Rice University

Achilles Benetopoulos  
(Unaffiliated)

Lazar Cvetković  
University of Belgrade

✉ [pash-discuss@googlegroups.com](mailto:pash-discuss@googlegroups.com)  
🐙 [github.com/andromeda/pash](https://github.com/andromeda/pash)

\* equal contribution

# Shell Scripts are Everywhere

Default/scriptable system interface  
even in the lightest containers  
Kubernetes, Docker

```
# Check all possible clusters, as your .kubeconfig may have multiple contexts:
kubectl config view -o jsonpath='{{"Cluster name\tServer\n"}{range .clusters[*]}{'

# Select name of cluster you want to interact with from above output:
export CLUSTER_NAME="gke_ps-dev-201405_us-east1_acaternberg"
#export SERVICE_ACCOUNT=cjoc
#export SERVICE_ACCOUNT=cloudbees-core-nginx-ingress
export SERVICE_ACCOUNT=default

# Point to the API server referring the cluster name
APISERVER=$(kubectl config view -o jsonpath="{.clusters[?(@.name==\"$CLUSTER_NAME)]}{'
```

Universal composition environment  
Commands (programs) can be written in  
C, C++, Rust, JS, Python, Ruby, Haskell...

```
echo "Building parser..."
cd compiler/parser
echo "|-- making libdash..."
make libdash &> $LOG_DIR/make_libdash.log
```

Succinct data processing:  
download/extraction/  
preprocessing/querying

```
seq $FROM $TO | sed "s;^;$IN;" | sed 's;$;/;' | xargs -r -n 1 $fetch | grep gz |
tr -s ' \n' | cut -d ' ' -f9 | sed 's;^(\.*)\((20[0-9][0-9]\)\.gz;/2/1\2\.gz;' |
sed "s;^;$IN;" | xargs -n1 curl -s | gunzip | # end-preprocessing
cut -c 89-92 | grep -v 999 | sort -rn | head -n1 # actual processing
```

## Bentley: A word-counting challenge

```
10 was
10 the
10 of
10 it
  2 times
```

```
-cs A-Za-z '\n' | tr A-Z a-z | sort | uniq -c | sort -rn | sed ${1}q
```

[illegible]

```
tr -cs A-Za-z '\n' | tr A-Z a-z | sort | uniq -c | sort -rn | sed ${1}q
```

It was the best of times, it was the  
worst of times, it was the age of  
wisdom, it was the age of  
foolishness, it was the epoch of  
belief, it was the epoch of  
incredulity, it was the season of  
Light, it was the season of  
Darkness, it was the spring of  
hope, it was the winter of despair.



```
tr -cs A-Za-z '\n'
```

It  
was  
the  
best  
of  
times  
it  
was  
the  
...

```
tr -cs A-Za-z '\n' | tr A-Z a-z | sort | uniq -c | sort -rn | sed ${1}q
```

It  
was  
the  
best  
of  
times  
it  
was  
the  
...

**tr A-Z a-z**

it  
was  
the  
best  
of  
times  
it  
was  
the  
...

tr -cs A-Za-z '\n' | **tr A-Z a-z** | sort | uniq -c | sort -rn | sed \${1}q

it  
was  
the  
best  
of  
times  
it  
was  
the  
...



sort

age  
age  
belief  
best  
darkness  
despair  
epoch  
epoch  
foolishness  
...

```
tr -cs A-Za-z '\n' | tr A-Z a-z | sort | uniq -c | sort -rn | sed ${1}q
```

```
age
age
belief
best
darkness
despair
epoch
epoch
foolishness
...
```

**uniq -c**

```
2 age
1 belief
1 best
1 darkness
1 despair
2 epoch
1 foolishness
1 hope
10 it
...
```

```
tr -cs A-Za-z '\n' | tr A-Z a-z | sort | uniq -c | sort -rn | sed ${1}q
```



```
2 age
1 belief
1 best
1 darkness
1 despair
2 epoch
1 foolishness
1 hope
1 incredulity
10 it
```

**sort -rn**

```
10 was
10 the
10 of
10 it
2 times
2 season
2 epoch
2 age
1 worst
1 wisdom
```

```
tr -cs A-Za-z '\n' | tr A-Z a-z | sort | uniq -c | sort -rn | sed ${1}q
```

```
10 was
10 the
10 of
10 it
 2 times
 2 season
 2 epoch
 2 age
 1 worst
 1 wisdom
...
```



`sed ${1}q`

```
10 was
10 the
10 of
10 it
 2 times
```

```
tr -cs A-Za-z '\n' | tr A-Z a-z | sort | uniq -c | sort -rn | sed ${1}q
```

It was the best of times, it was the  
worst of times, it was the age of  
wisdom, it was the age of  
foolishness, it was the epoch of  
belief, it was the epoch of  
...

It was the best of times, it was the  
worst of times, it was the age of  
wisdom, it was the age of  
foolishness, it was the epoch of  
belief, it was the epoch of  
...

It was the best of times, it was the  
worst of times, it was the age of  
wisdom, it was the age of  
foolishness, it was the epoch of  
belief, it was the epoch of  
...

It was the best of times, it was the  
worst of times, it was the age of  
wisdom, it was the age of  
foolishness, it was the epoch of  
belief, it was the epoch of  
...



10 was  
10 the  
10 of  
10 it  
2 times

# How to parallelize?

```
tr -cs A-Za-z '\n' | tr A-Z a-z | sort | uniq -c | sort -rn | sed ${1}q
```

# Shell scripts are mostly sequential

Their parallelization requires considerable effort:

- Command-specific flags (e.g., **sort -p**, **make -jN**)
- Mostly-manual, restricted parallelization tools (e.g., GNU **parallel**)
- Full rewrites in parallel frameworks (e.g., MapReduce)

```

import java.io.*;
import java.util.*;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.mapreduce.Mapper;

public class top_10_Movies_Mapper extends Mapper<Object,
Text, Text, LongWritable> {

    private TreeMap<Long, String> tmap;

    @Override
    public void setup(Context context) throws IOException,
    InterruptedException
    {
        tmap = new TreeMap<Long, String>();
    }

    @Override
    public void map(Object key, Text value,
        Context context) throws IOException,
    InterruptedException
    {
        // no_of_views (tab separated)
        // we split the input data
        String[] tokens = value.toString().split("\t");

        String movie_name = tokens[0];
        long no_of_views = Long.parseLong(tokens[1]);

        tmap.put(no_of_views, movie_name);

        if (tmap.size() > 10)
        {
            tmap.remove(tmap.firstKey());
        }
    }

    @Override
    public void cleanup(Context context) throws IOException,
    InterruptedException
    {
        for (Map.Entry<Long, String> entry : tmap.entrySet())
        {
            long count = entry.getKey();
            String name = entry.getValue();
            context.write(new Text(name), new LongWritable(count));
        }
    }
}

import java.io.IOException;
import java.util.Map;
import java.util.TreeMap;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class top_10_Movies_Reducer extends Reducer<Text,
LongWritable, LongWritable, Text> {
    private TreeMap<Long, String> tmap2;

    @Override
    public void setup(Context context) throws IOException,
    InterruptedException
    {
        tmap2 = new TreeMap<Long, String>();
    }

    @Override
    public void reduce(Text key, Iterable<LongWritable> values,
        Context context) throws IOException, InterruptedException
    {
        String name = key.toString();
        long count = 0;

        for (LongWritable val : values) {
            count = val.get();
        }

        tmap2.put(count, name);

        if (tmap2.size() > 10) {
            tmap2.remove(tmap2.firstKey());
        }
    }

    @Override
    public void cleanup(Context context) throws IOException,
    InterruptedException {
        for (Map.Entry<Long, String> entry : tmap2.entrySet()) {
            long count = entry.getKey();
            String name = entry.getValue();
            context.write(new LongWritable(count), new Text(name));
        }
    }
}

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;

public class Driver {
    public static void main(String[] args) throws Exception
    {
        Configuration conf = new Configuration();
        String[] otherArgs = new GenericOptionsParser(conf,
args).getRemainingArgs();

        if (otherArgs.length < 2)
        {
            System.err.println("Error: please provide two paths");
            System.exit(2);
        }

        Job job = Job.getInstance(conf, "top 10");
        job.setJarByClass(Driver.class);

        job.setMapperClass(top_10_Movies_Mapper.class);
        job.setReducerClass(top_10_Movies_Reducer.class);

        job.setMapOutputKeyClass(Text.class);
        job.setMapOutputValueClass(LongWritable.class);

        job.setOutputKeyClass(LongWritable.class);
        job.setOutputValueClass(Text.class);

        FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
        FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```

# Big-Data Version of McIlroy's Pipeline

## 150-line Hadoop Program

# Mostly sequential by default — how to parallelize?

Parallelization requires considerable effort:

- Command-specific flags (e.g., `sort -p`, `make -jN`)
- Mostly-manual, restricted parallelization tools (e.g., GNU `parallel`)
- Full rewrites in parallel frameworks (e.g., MapReduce)

# Challenges of Automating Shell-Script Parallelization

```
for directory in /project/gutenberg/*/; do  
  ls $directory | grep 'txt' | wc -l > index.txt  
done
```

```
cat f1 f2 |  
tr -cs A-Za-z '\n' | tr A-Z a-z | sort | uniq -c | sort -rn | sed ${1}q
```

split

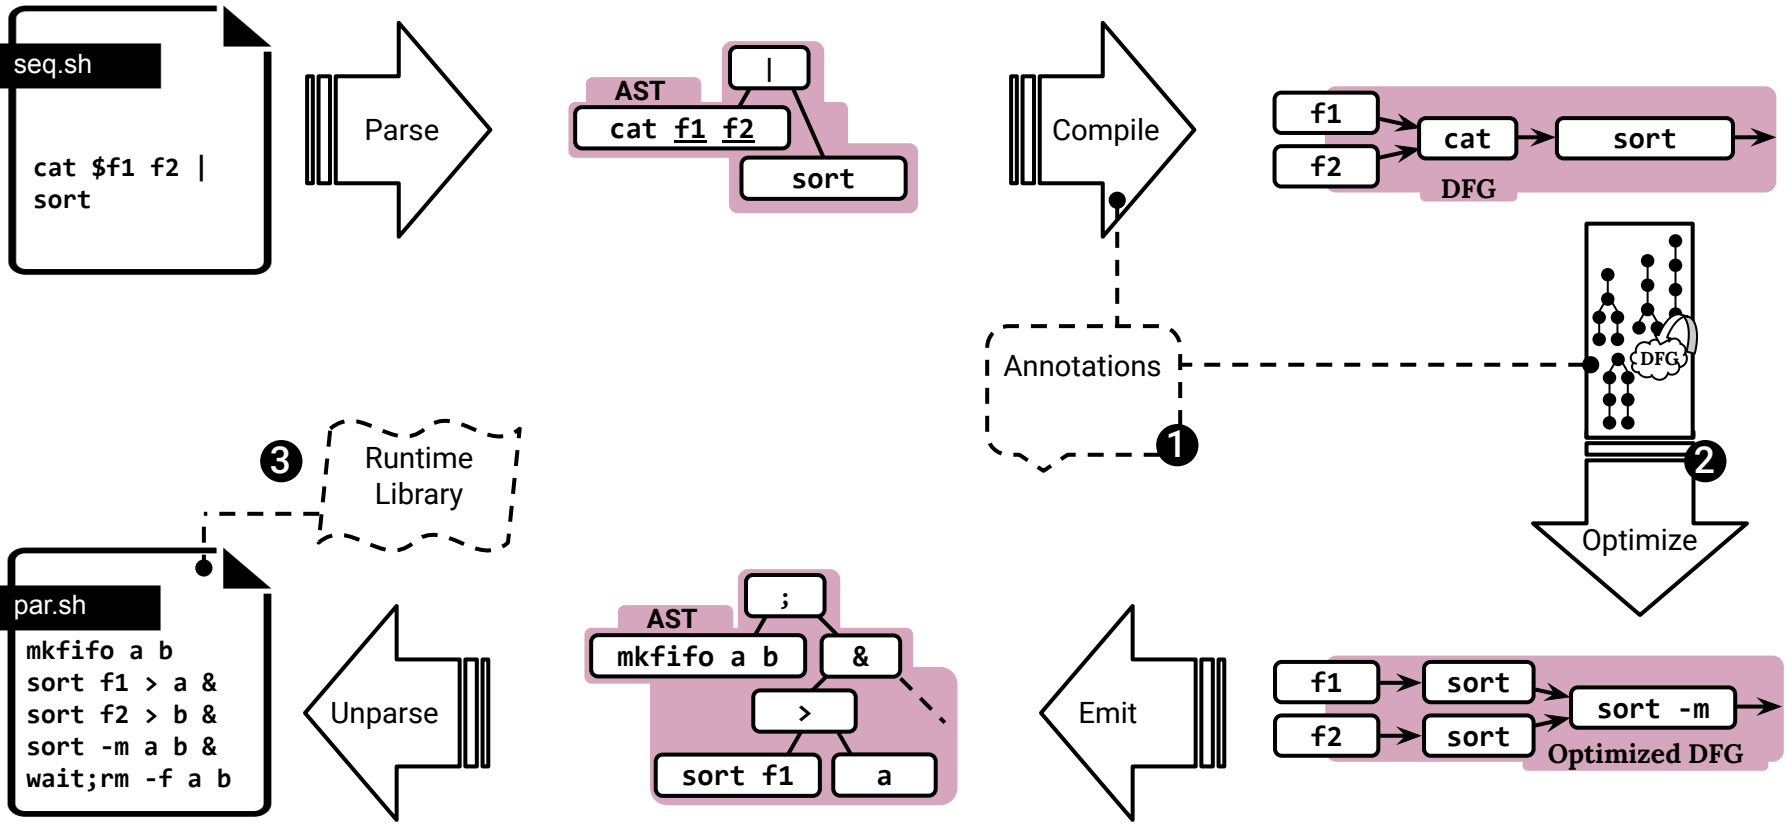


aggregate

```
echo 'Done';
```


- (1) Numerous and opaque Unix commands
- (2) Shell language enforced dependencies
- (3) Runtime support for Unix parallelization

# PaSh Overview






# PaSh Overview



# 1. Unix Parallelizability Study & Annotations

# POSIX GNU



Parallelizability properties:

- 4 broad classes
- Flags and options
- Input consumption

Parallelizability DSL:

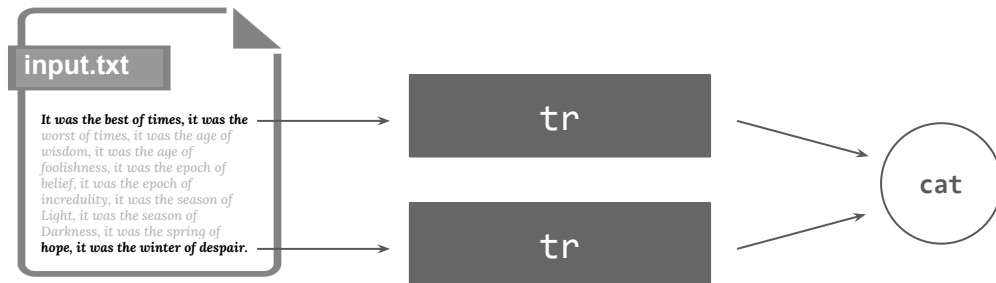
(cmd,  
flg,  
[in]) → DFG node

4

command  
parallelizability  
classes

12.7%

stateless



4


command  
parallelizability  
classes

12.7%

stateless

8.7%

parallelizable pure



# 4

command  
parallelizability  
classes

12.7%

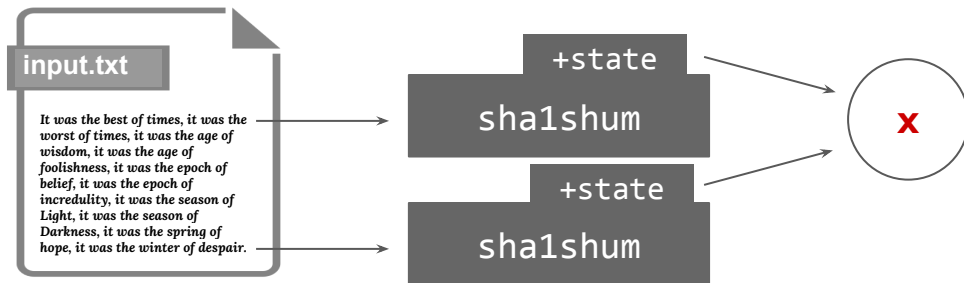
stateless

8.7%

parallelizable pure

8.2%


non-parallelizable pure



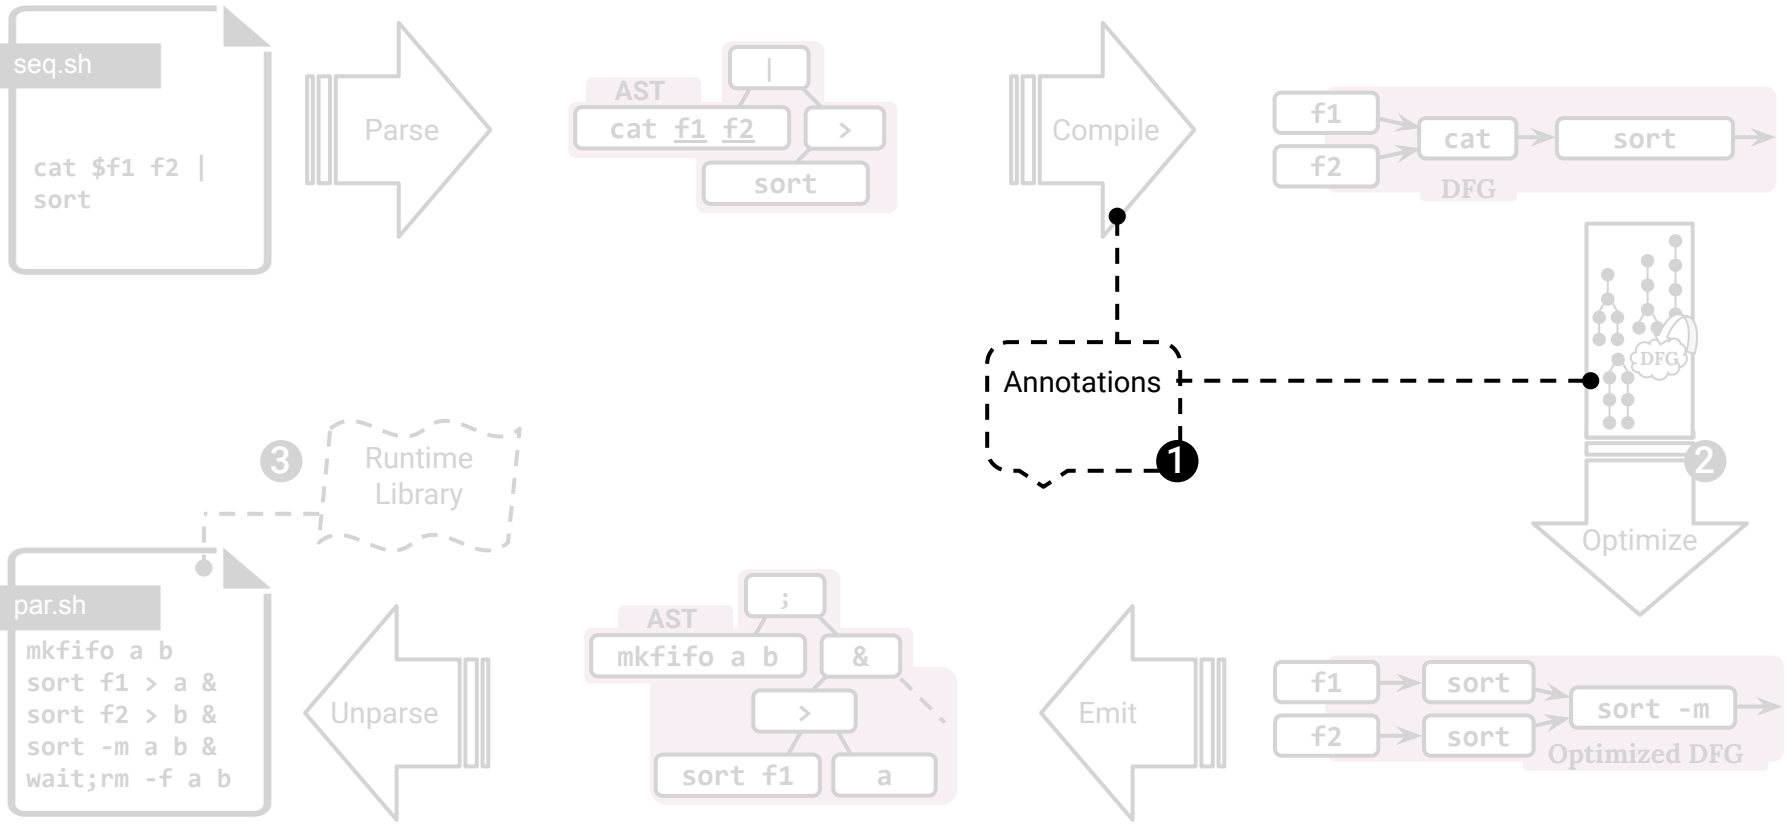
# 4

command  
parallelizability  
classes

12.7%	stateless
8.7%	parallelizable pure
8.2%	non-parallelizable pure
70.4%	side-effectful

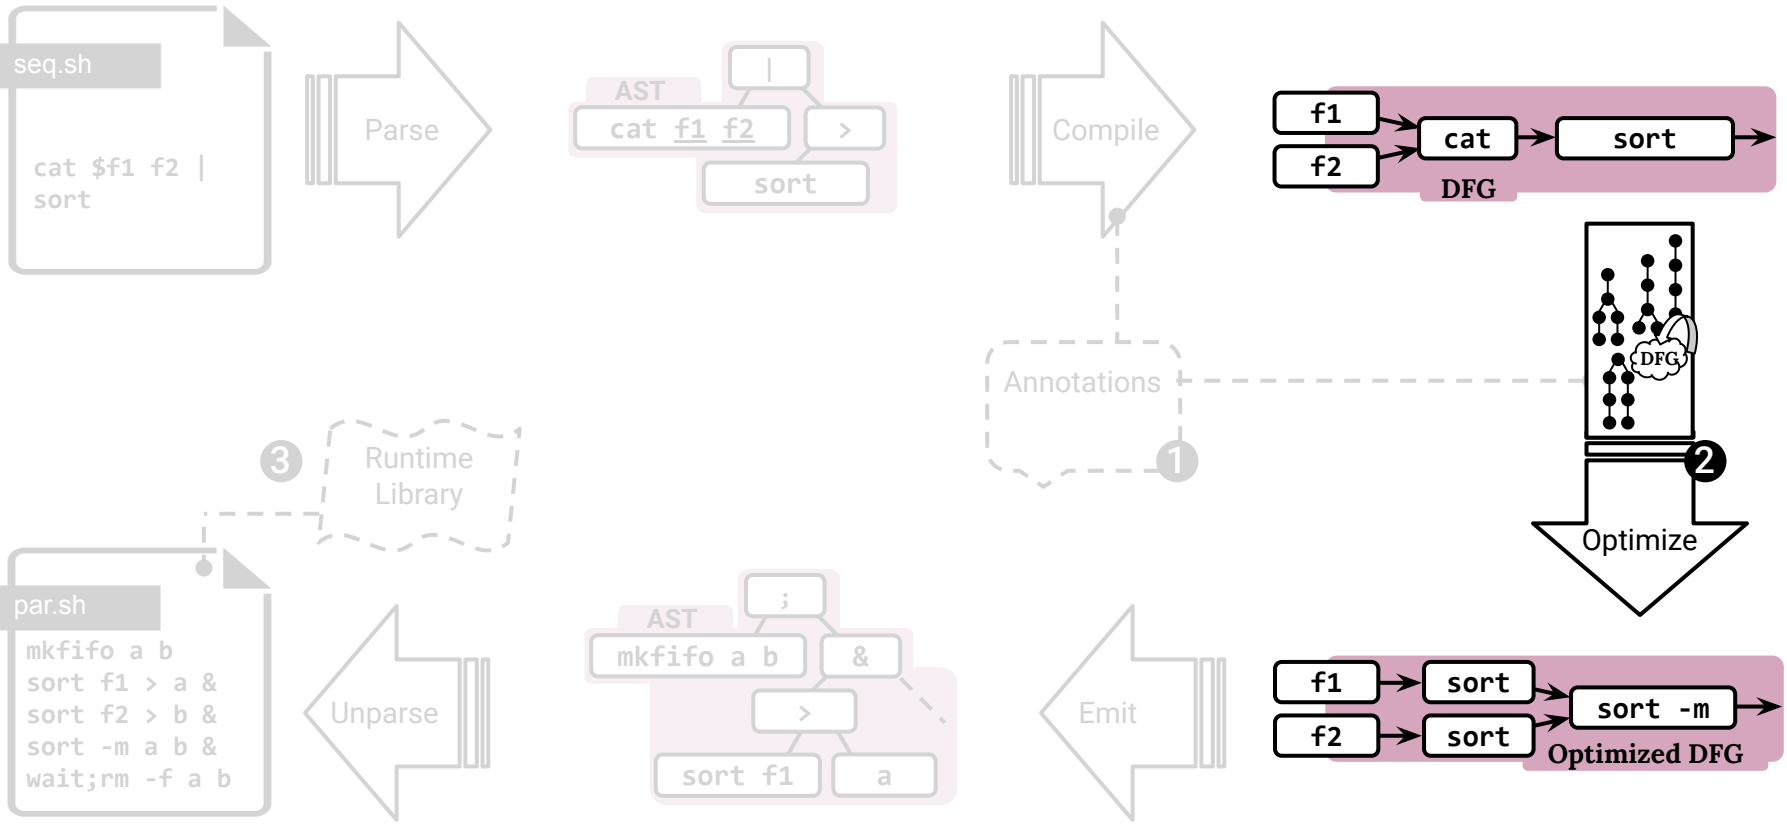


# PaSh Overview






# PaSh Overview




## 2. Dataflow Model & Transformations

*Scheduling constraint*


```
cat f1 f2 > out.txt; cat out.txt
```



```
cat f1 f2 | tr A-Z a-z | sort > out.txt
```




```
cat f1 f2 | tr A-Z a-z | sort > out.txt
```




Transformation condition: `tr` is stateless

```
cat f1 f2 | tr A-Z a-z | sort > out.txt
```




Transformation condition: `cat` followed by `split`

```
cat f1 f2 | tr A-Z a-z | sort > out.txt
```



Transformation condition: **sort** is parallelizable pure


```
cat f1 f2 | tr A-Z a-z | sort > out.txt
```



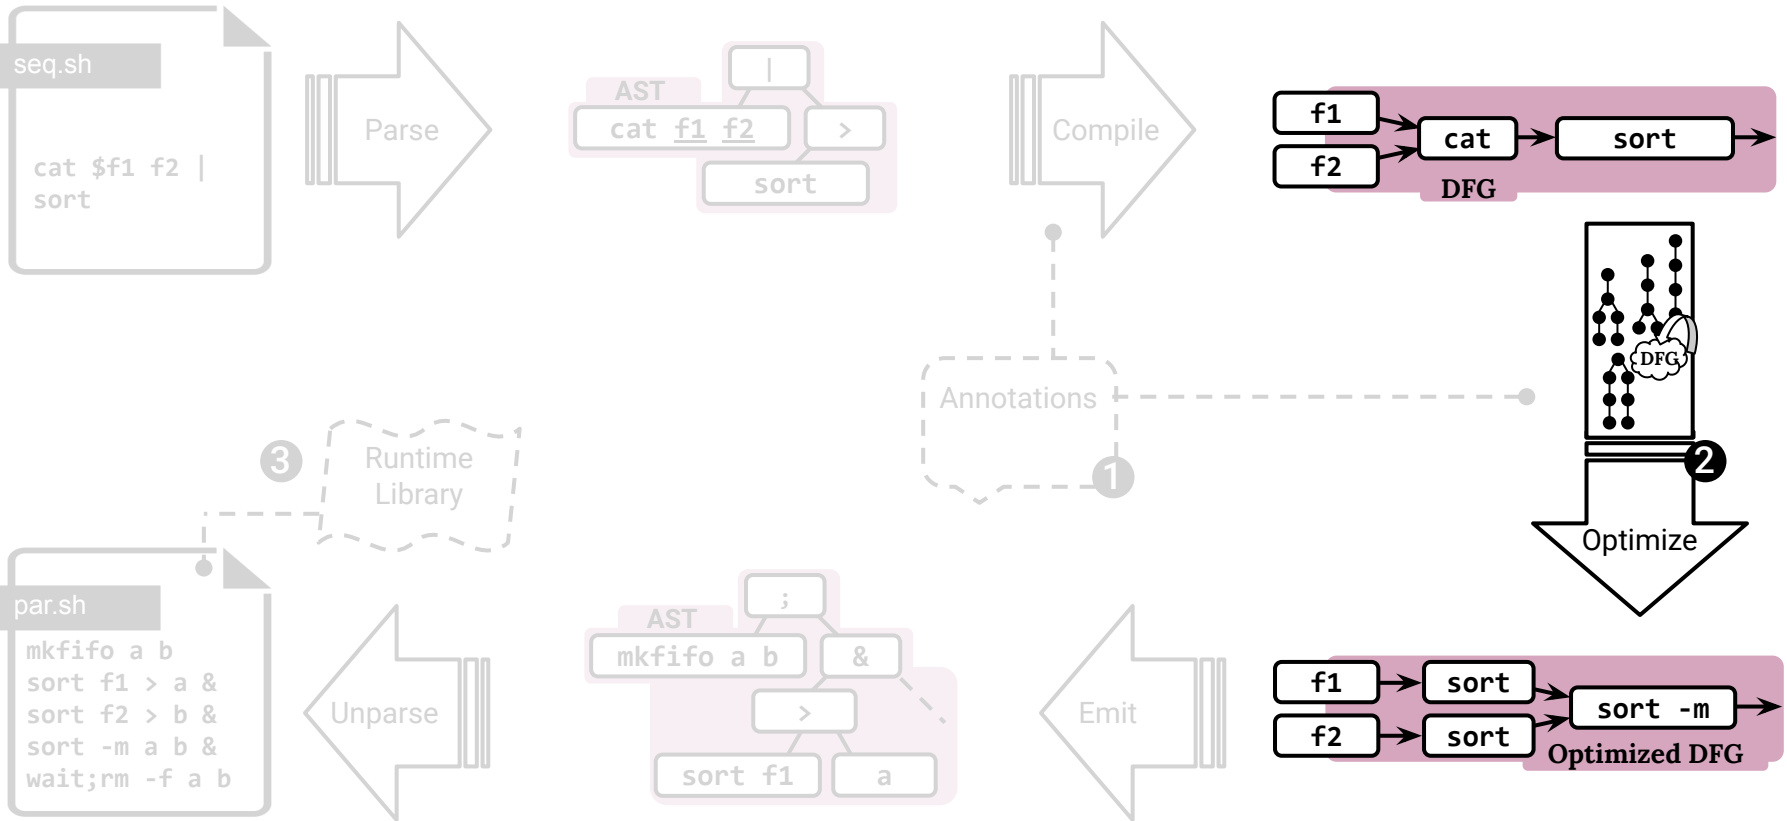
Transformation condition: `cat` followed by `split`



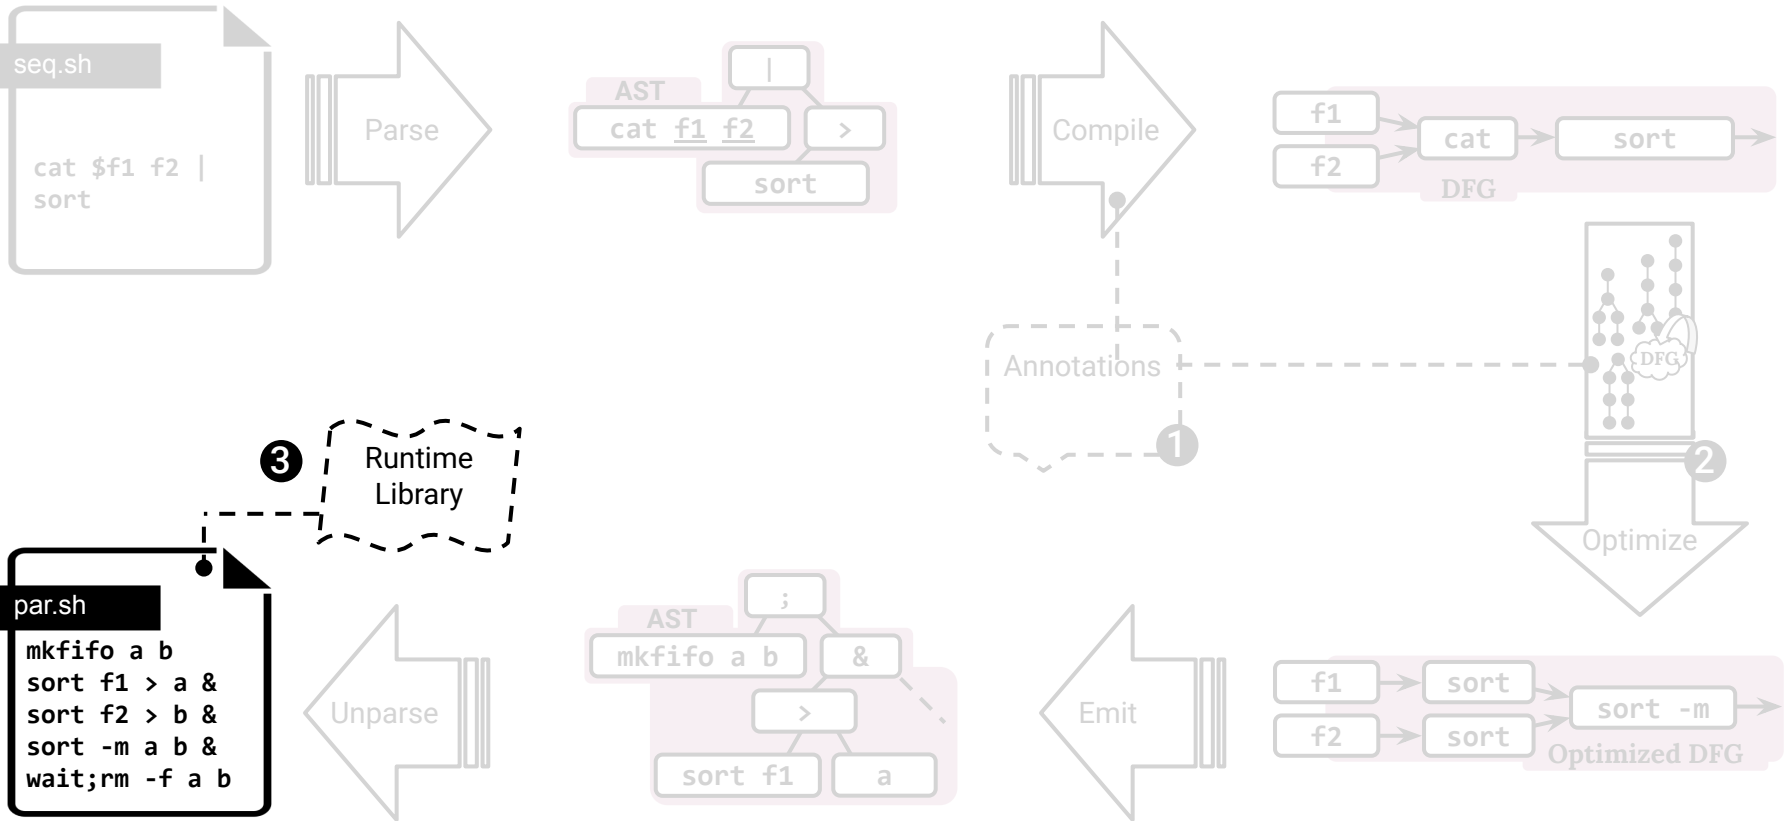
# 1 + 3 Transformations



# PaSh Overview



# PaSh Overview




### 3. Runtime Support

# Runtime Support: Performance & Correctness

- Unix pipes are lazy, *i.e.*, inadequate buffering (and for a good reason)
- Dataflow graph termination is tricky
- *Parallelizable-pure* commands require careful aggregation


# Runtime Challenge: Unix's Lazy Semantics

```
mkfifo f1 f2  
grep "foo" in1 > f1 &  
grep "foo" in2 > f2 &  
cat f1 f2
```




# Runtime Challenge: Unix's Lazy Semantics

```
mkfifo f1 f2  
grep "foo" in1 > f1 &  
grep "foo" in2 > f2 &  
cat f1 f2
```



# Runtime Challenge: Unix's Lazy Semantics


```
mkfifo f1 f2  
grep "foo" in1 > f1 &  
grep "foo" in2 > f2 &  
cat f1 f2
```





# Runtime Challenge: Unix's Lazy Semantics

```
mkfifo f1 f2  
grep "foo" in1 > f1 &  
grep "foo" in2 > f2 &  
cat f1 f2
```




# Runtime Challenge: Unix's Lazy Semantics

```
mkfifo f1 f2
```

```
grep "foo" in1 > f1 &
```


```
grep "foo" in2 > f2 &
```

```
cat f1 f2
```



# Runtime Challenge: Unix's Lazy Semantics



```
mkfifo f1 f2  
grep "foo" in1 > f1 &  
grep "foo" in2 > f2 &  
cat f1 f2
```



Execution proceeds in steps!

# A non-solution: Use intermediary files...


```
touch f1 f2
grep "foo" in1 > f1 &
grep "foo" in2 > f1 &
wait
cat f1 f1
```



Among other problems, this "solution" prevents pipeline parallelism (more on that later)


# The PaSh Solution: Eager Buffers

```
mkfifo f1 f2 f3 f4
grep "foo" in1 > f1 &
grep "foo" in2 > f2 &
eager < f1 > f3 &
eager < f2 > f4 &
cat f3 f4
```




# The PaSh Solution: Eager Buffers

```
mkfifo f1 f2 f3 f4
grep "foo" in1 > f1 &
grep "foo" in2 > f2 &
eager < f1 > f3 &
eager < f2 > f4 &
cat f3 f4
```




# The PaSh Solution: Eager Buffers

```
mkfifo f1 f2 f3 f4
grep "foo" in1 > f1 &
grep "foo" in2 > f2 &
eager < f1 > f3 &
eager < f2 > f4 &
cat f3 f4
```



# The PaSh Solution: Eager Buffers


```
mkfifo f1 f2 f3 f4
grep "foo" in1 > f1 &
grep "foo" in2 > f2 &
eager < f1 > f3 &
eager < f2 > f4 &
cat f3 f4
```





# The PaSh Solution: Eager Buffers

```
mkfifo f1 f2 f3 f4
grep "foo" in1 > f1 &
grep "foo" in2 > f2 &
eager < f1 > f3 &
eager < f2 > f4 &
cat f3 f4
```



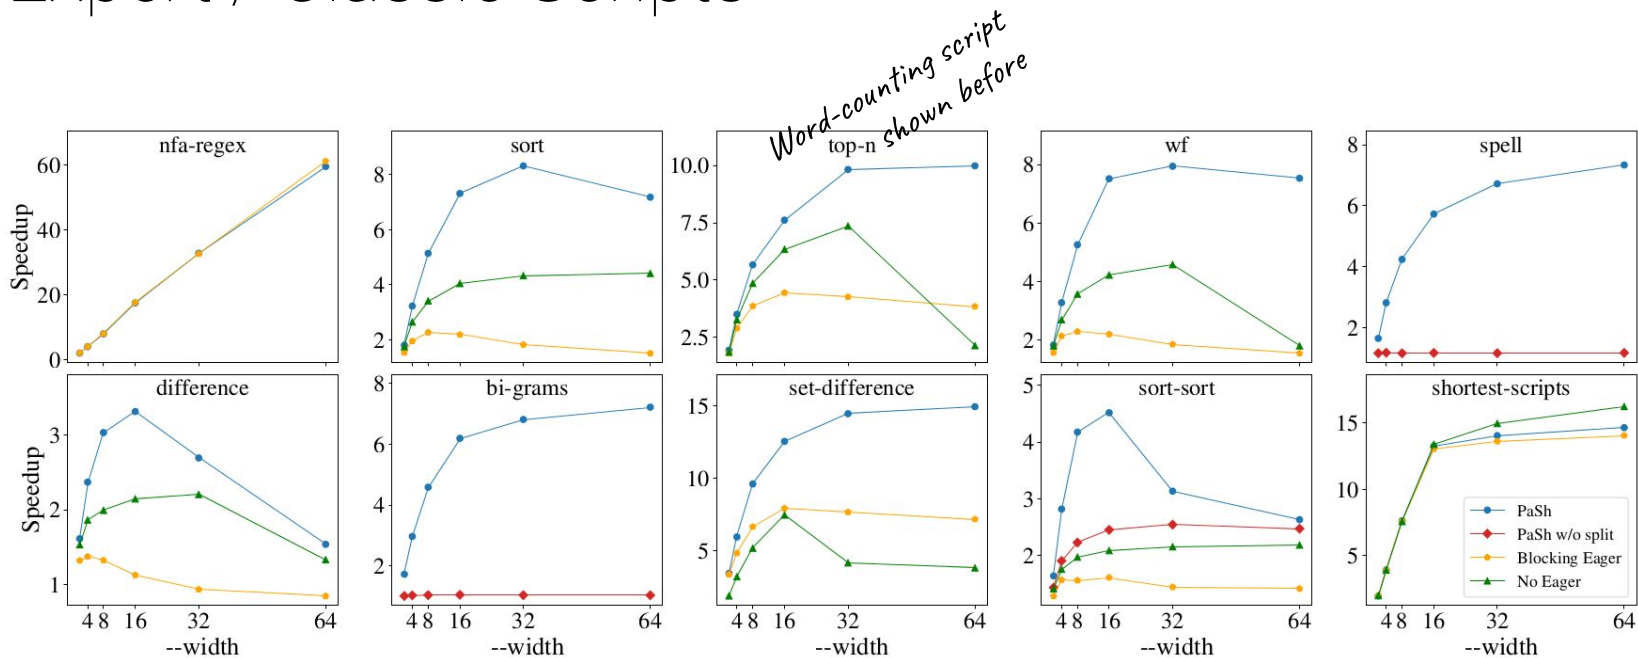
/pash/runtime/eager

- Unix command, usable outside PaSh too
- Buffers input eagerly — can spill to disk
- Keeps fragment in DFG model

Demo Time!


Evaluation

# 1. Expert / Classic Scripts




Speedups against bash baseline  
for `pash --width=16`:

5.93x vs. 8.83x



## 2. Pipelines in the wild



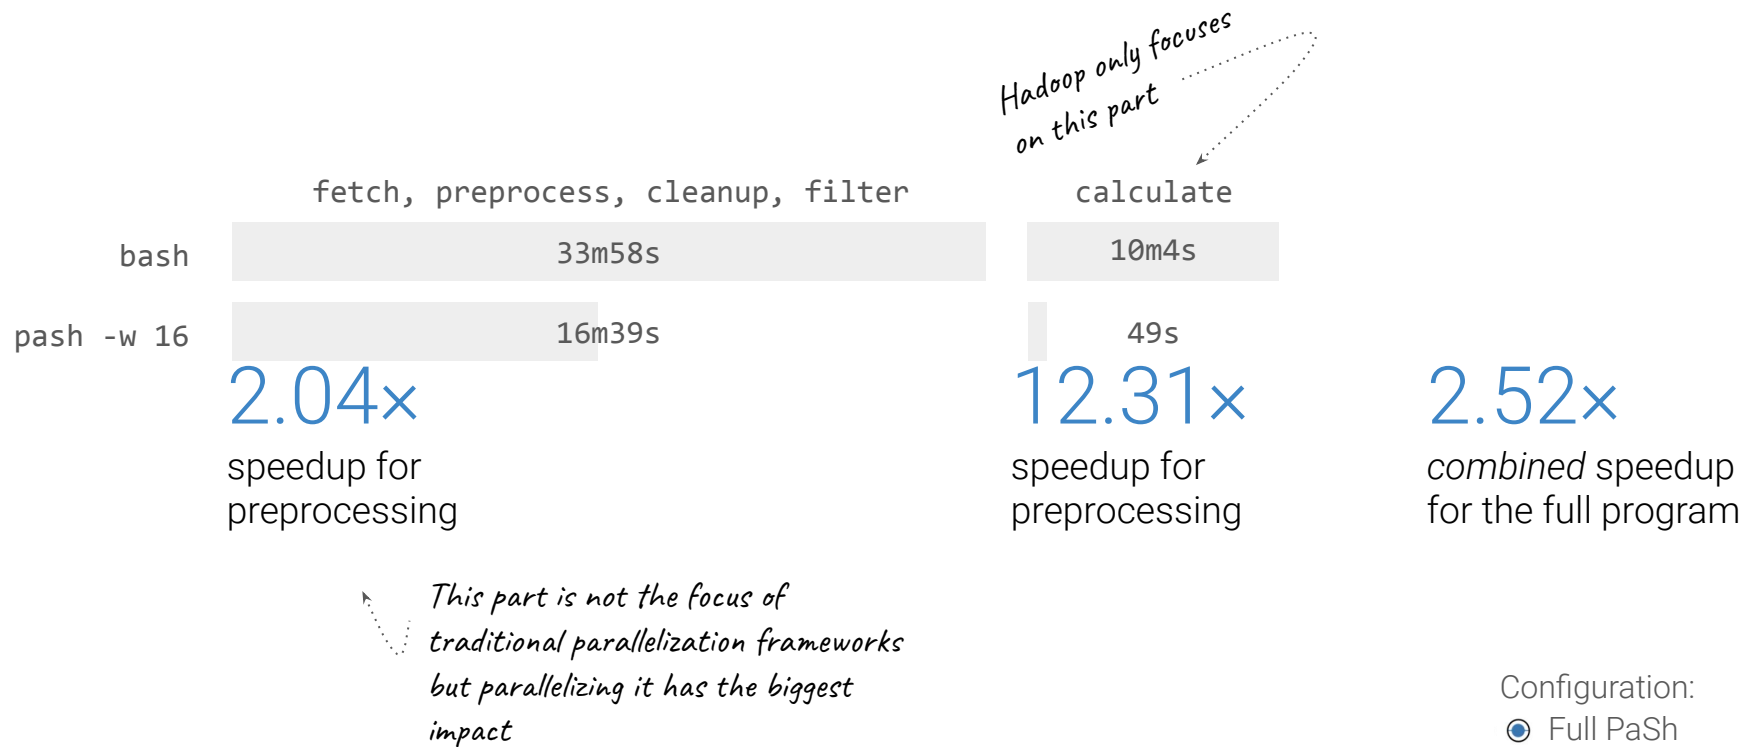
+ PaSh awareness goes a long way!

e.g. #26 { `cat $IN6 | awk '{print $2, $0}' | sort -nr | cut -d ' ' -f 2` (1.01x)  
`cat $IN6 | sort -nr -k2 | cut -d ' ' -f 1` (8.1x !!1!1)

Configuration:

● Full PaSh  
 --width=16

### 3. Case Study no.1: NOAA Weather Analysis



Configuration:  
● Full PaSh  
--width=16  
82GB (5y data)

Conclusion

# Conclusion

- Parallelize unix shell scripts (POSIX -> POSIX)
- Annotations address extensibility issues
- Open source — 12+ contributors
- Lots of recent excitement — let's rehabilitate the shell!

